

Chapter 1

INTRODUCTION

1.1 BACKGROUND

1.1.1 Estimation and Related Areas

Estimation is the process of inferring the value of a quantity of interest from indirect, inaccurate and uncertain observations.

The purpose of estimation can be:

- Determination of planet orbit parameters — probably the first estimation problem — studied by Laplace, Legendre, and Gauss
- Statistical inference
- Determination of the position and velocity of an aircraft in an air traffic control system — tracking
- Application of control to a plant in the presence of uncertainty (noise and/or unknown parameters) — parameter identification, state estimation and stochastic control
- Determination of model parameters for predicting the state of a physical system or forecasting economic or other variables — system identification
- Determination of the characteristics of a transmitted message from noise-corrupted observation of the received signal — communication theory
- Determination of some parameters or characteristics of a signal or image — signal/image processing

More rigorously, *estimation* can be viewed as the *process of selecting a point from a continuous space* — the “best estimate.”

Decision can be viewed as the *selection of one out of a set of discrete alternatives* — the “best choice” from a discrete space. However, one can talk about estimation in a discrete-valued case with the possibility of not making a choice but obtaining some conditional probabilities of the various alternatives. This information can be used without making “hard decisions.”

2 1 INTRODUCTION

Estimation and decision can, therefore, be seen to be overlapping and techniques from both areas are used simultaneously in many practical problems.

Tracking is the estimation of the state of a moving object based on *remote measurements*. This is done using one or more sensors at fixed locations or on moving platforms.

At first sight, tracking might seem to be a special case of estimation. However, it is wider in scope: Not only does it use all the tools from estimation, but it also requires extensive use of statistical decision theory when some of the practical problems (data association — “Which is my measurement?” — see, e.g., [Bar-Shalom95]) are considered.

Filtering is the estimation of the (current) state of a dynamic system — the reason for the use of the word “filter” is that the process for obtaining the “best estimate” from noisy data amounts to “filtering out” the noise. The term filtering is thus used in the sense of eliminating an undesired signal, which, in this case, is the noise.

In control systems, signals are also filtered to obtain the estimate of the state of the (noisy) dynamic system, needed by the controller. Filtering of signals is very commonly used in signal processing — in the frequency domain as well as in the spatial domain. The latter is done, for example, to select signals coming from a certain direction. **Navigation** is the estimation of the state of the platform (“own ship”) on which the sensor is located.

An **optimal estimator** is a computational algorithm that processes observations (measurements) to yield an estimate of a variable of interest, which optimizes a certain criterion.

In view of the above discussion, estimation can be viewed as a scheme for **information extraction and enhancement**: Based on measurements (observations), we want to maximize our knowledge about a parameter, a state, a signal, an image, and so on.

In general, one can classify the variable that is to be estimated into the following two categories:

- A parameter — a time-invariant quantity (a scalar, a vector, or a matrix)
- The state of a dynamic system (usually a vector), which evolves in time according to a stochastic equation

Consequently, one has two classes of estimators:

- Parameter estimators
- State estimators

While some engineering texts deal mainly (or exclusively) with state estimation, which is more general than parameter estimation, a balanced approach is taken here. The reason it is important to understand both is that some results — e.g., convergence and, more importantly, **information limits** — are more readily available/usable for parameter estimators than for state estimators. This

understanding can be helpful in choosing the type of mathematical model used for solving practical problems.

Figure 1.1.1-1 presents a concise block diagram that illustrates state estimation. In this figure the first two blocks are “black boxes” — there is no access to variables inside them. The only variables to which the estimator has access are the **measurements**, which are affected by the error sources in the form of “measurement noise.”

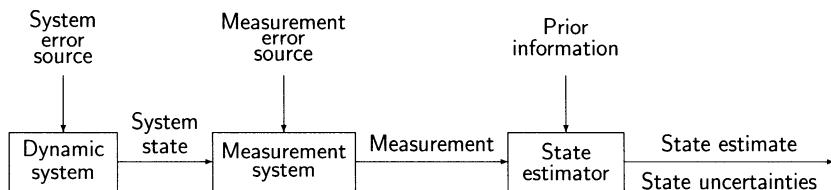


Figure 1.1.1-1: State estimation — information extraction and enhancement.

The estimator uses knowledge about the following:

- The evolution of the variable (the system dynamics)
- The measurement model
- The probabilistic characterization of the various random factors (disturbances)
- The prior information

Optimal Estimation — Advantages and Disadvantages

The advantage of an optimal estimator is that it makes the best utilization of the data and the knowledge of the system and the disturbances.

The disadvantages, like for any optimal technique, are that it is possibly sensitive to modeling errors and might be computationally expensive. In view of this, it is very important to have a clear understanding of the assumptions under which an algorithm is optimal and how they relate to the real world.

1.1.2 Applications of Estimation

State estimation, in combination with decision theory, has a great variety of applications. A partial list of the areas where this has found use is as follows:

- Tracking/surveillance and trajectory determination
- Navigation
- Control systems
 - guidance
 - attitude control

- sensor pointing
- steel making
- chemical, nuclear, and industrial processes, etc.
- Power systems
- Failure detection
- Signal processing
- Image processing
- Communications
- Biomedical engineering
- Operations research
- Mapping via remote sensing
- Geophysical problems
- Fluid flow rate measurement
- Econometric systems
 - macroeconomic models
 - microeconomic models
- Demographic systems

1.1.3 Preview of Estimation/Filtering

This subsection will present a bird's-eye view of estimation, from the the Least Squares technique originated by Gauss and Legendre for parameter estimation in astronomical studies (motion of heavenly bodies) to Kalman filtering for state estimation. All this will be done without mathematics, to illustrate the concepts.¹

Estimation is the process of inferring the values of some parameters that characterize, e.g., the motion of a heavenly body (the six parameters of planetary motion — 3 position and 3 velocity — evolving in a predictable manner) or the state of an object, e.g., an aircraft (evolving in a not completely predictable manner).

Gauss made the following (philosophical) observations on the (physical) observations that could be made on the planetary motion:

- If the observations were absolutely correct, the parameters could be determined with perfect accuracy from a minimum number of observations (n observations for n parameters).
- Then subsequent observations would confirm, but not correct, the values obtained for the parameters.

¹The alternative title for this subsection could have been "Kalman Filtering for Dummies." This material could be used by readers eager to show their spouses how they make a living; however, we assume no responsibility for the possible ensuing marital problems.

- But, since the observations are only approximations of the truth, we should combine more observations than the minimum to determine more accurately the unknown quantities.
- Thus, starting with approximate knowledge, we can correct it with subsequent observations so as to satisfy all the observations in the most accurate manner possible.

The above (philosophical) observations lead to the following implications:

- A basic description (model) of the “system” is available with some unknown parameters to be estimated.
- Redundant data are required to reduce the effect of measurement errors.
- To satisfy all the observations in the most accurate manner possible, the “residuals” (differences between the observed values and the values predicted from the estimates) should be as small as possible.
- Inaccuracy of observations necessitates probabilistic modeling.
- Combination of the initial knowledge and the subsequent observations leads to the recursive algorithm concept.

Estimation theory provides a systematic approach to deal with the discrepancies in different measurements on the same object.

There are a number of estimation problems that will be briefly outlined below.

The Least Squares (Gauss-Legendre) Problem

If the observations consist of

- linear combinations of the parameters of interest, in the presence of additive errors (noise) assumed to be zero-mean Gaussian random variables,

then the

- minimization of the sum of the squares of the residuals (errors) in fitting the data namely, the least squares (LS) of the difference between the model and the observations,

is equivalent to the

- maximization of the likelihood function of the parameters (the probability density function of the measurements conditioned on the unknown parameters).

The Wiener-Hopf Problem

The variables to be estimated are not constant but time-varying: The objective is to estimate a random process.

The process to be estimated, $x(t)$, is assumed stationary (typically, zero-mean) and with known autocorrelation. The cross-correlation between this process and the observations $z(t)$ is assumed also known.

The solution that yields the estimate $\hat{x}(t)$ of the process, obtained in the frequency domain as a transfer function $W_o(s)$ via Fourier transforms, so as to minimize a mean square error is the Wiener filter. This is illustrated in Fig. 1.1.3-1.

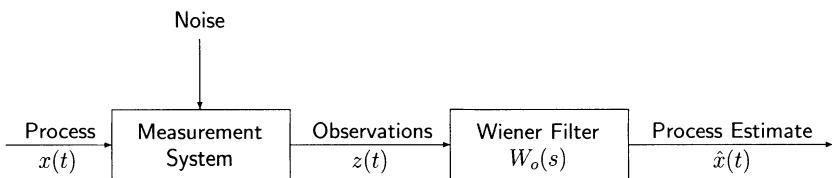


Figure 1.1.3-1: The Wiener filter — linear time-invariant system with transfer function $W_o(s)$.

The Kalman-Bucy Problem

The object of the estimation is a random process, which is not necessarily stationary. Its characterization is given by a linear differential (or difference) equation with known coefficients — the state equation of a dynamic system. Random disturbances affect the evolution of the process (process noise). The process is observed in the presence of measurement errors (measurement noise).

The solution obtained (in the time domain) so as to minimize a mean square error is the Kalman-Bucy or Kalman filter.

Solution to the Least Squares Estimation Problem

The standard (batch) solution to the Least Squares Estimation (LSE) problem consists of the following:

- Given k measurements (each an m -vector) of a constant parameter (an n -vector), the LS estimate of the unknown parameter is obtained by processing simultaneously the entire measurement set (a km -vector).
- To obtain the new estimate of the parameter when one gets a new measurement (the $(k+1)$ th m -vector), one reprocesses the entire data, a $(k+1)m$ -vector, i.e., a new problem of *larger size* is solved.

The batch LSE solution can be rewritten (if the noise components in the k sets of measurements are uncorrelated) as a *recursion*

$$(\text{estimate})_{k+1} = (\text{estimate})_k + (\text{weighting})_{k+1} \times (\text{residual})_{k+1}$$

where k is the time index, $(\cdot)_k$ indicates the value at time k and

$$\begin{aligned} (\text{residual})_{k+1} &= (\text{measurement})_{k+1} \\ &\quad - (\text{predicted measurement})_{k+1} \text{ based on } (\text{estimate})_k \end{aligned}$$

The size of this problem is *the same* for every k .

The accuracy of the estimate also obeys a recursion; that is, the “precision” of the estimate (the inverse of the parameter mean square error — MSE) evolves according to

$$\begin{aligned} (\text{precision})_{k+1} &= (\text{precision})_k \\ &\quad + (\text{information about the parameter in measurement})_{k+1} \end{aligned}$$

Thus the “precision” increases (the MSE or variance decreases) strictly if the measurements contain information about the parameters — if the parameters are *observable*.

If the parameters are constant, as $k \rightarrow \infty$ one accumulates an unlimited amount of information and the error in the estimates decreases to zero — the LSE converges to the true value. As the MSE in the estimate decreases, so does the weighting in the recursion: as the confidence in the estimates grows, new data are weighted less and less.

Estimation of the State of a Linear Stochastic Dynamic System

A realistic model for the state of a system has typically

- an imperfect initial estimate
- an imperfectly predictable evolution (a stochastic disturbance affects its dynamics — the process noise)

Similarly, the measurement model is as follows:

- Only some state components (or combinations of them) are observed.
- These observations are inaccurate — corrupted by measurement noise.

The evolution model for the state (discrete time dynamic equation) is as follows:

$$(\text{state})_{k+1} = (\text{linear function of the state})_k + (\text{process noise})_k$$

where k is the time index and $(\cdot)_k$ indicates the value at time k .

The observation process (measurement equation) is

$$(\text{measurement})_k = (\text{linear function of the state})_k + (\text{measurement noise})_k$$

8 1 INTRODUCTION

The two noise sequences are assumed to have known means (typically zero) and variances (covariance matrices in the multidimensional case), uncorrelated in time and from each other.

The (discrete time) Kalman filter (KF) computes the best estimate of the current state (in the minimum mean square error sense) based on the measurements up to the current time, denoted as (state estimate)_k. The KF recursion is as follows:

$$(\text{state estimate})_{k+1} = (\text{predicted state})_{k+1^-} + (\text{weighting})_{k+1} \times (\text{residual})_{k+1}$$

where $k + 1^-$ denotes “just before time $k + 1$ ” and

$$(\text{residual})_{k+1} = (\text{measurement})_{k+1} - (\text{predicted measurement})_{k+1^-}$$

The LSE recursion is remarkably similar:

$$\begin{aligned} (\text{parameter estimate})_{k+1} &= (\text{parameter estimate})_k \\ &\quad + (\text{weighting})_{k+1} \times (\text{residual})_{k+1} \end{aligned}$$

The accuracy of the estimate (precision — inverse of the variance) is as follows:

$$\begin{aligned} (\text{precision})_{k+1} &= (\text{precision})_{k+1^-} \\ &\quad + (\text{information about the state in measurement})_{k+1} \end{aligned}$$

The filter gain reflects the relative accuracy of the predicted state vs. the new observation. The new (updated) state estimate is the optimal combination of

- the entire past data (with the predicted state being the sufficient statistic) and
- the latest measurement

The sequence of the operations in the KF is outlined in Fig. 1.1.3-2. The calculations involving the state estimate use the measurements and have to be done in real time. However, the calculations involving the variances (covariance matrices in the multidimensional case) are data-independent and can be done off-line (before the data are obtained). The variance of the estimates (their “precision”) can be precomputed using the statistical description (variances) of the noise components entering into the system and the initial state variance.

Illustration for Tracking a Moving Point

Figure 1.1.3-3 illustrates graphically the state estimation process for a point moving with a nearly constant velocity along the x axis. The state vector consists of position x and velocity \dot{x} , with estimates \hat{x} and $\hat{\dot{x}}$, respectively. The

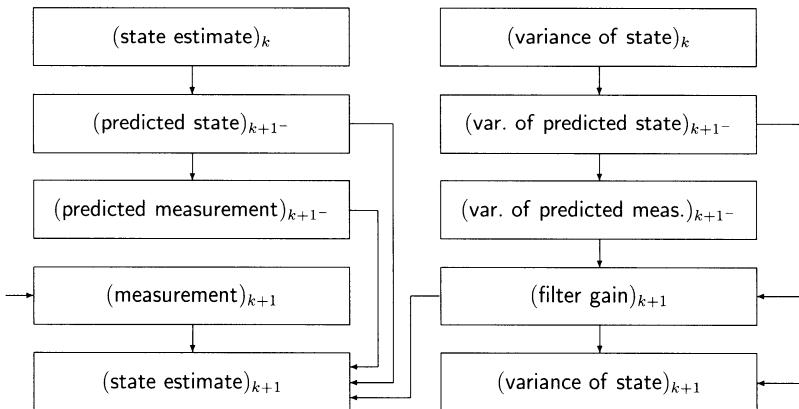


Figure 1.1.3-2: Sequence of operations in the KF.

covariance matrix of the state is denoted by P and it indicates the uncertainty region around the estimated state where the true state lies with high probability.

The measurement is

$$z_{k+1} = x_{k+1} + (\text{measurement noise})_{k+1}$$

The new position estimate will be between the observed and predicted values:

$$z_{k+1} < \hat{x}_{k+1} < \hat{x}_{k+1^-}$$

Assuming that at time $k+1$ the measured position is behind the predicted position, that is,

$$z_{k+1} - \hat{x}_{k+1^-} < 0$$

the new velocity estimate will be less than the predicted one:

$$\hat{\dot{x}}_{k+1} < \hat{\dot{x}}_{k+1^-}$$

Under stationarity (and state observability) conditions, the variance of the estimates converges to a steady-state value, as illustrated in Fig. 1.1.3-4, where the prediction and updated variances are shown.

Remarks

The KF is a stochastic counterpart of the (deterministic) state observer.

When the covariance is in steady state, the KF gain becomes a constant — such filters (for “nearly constant velocity” motion models) are known as α - β filters.

The Wiener filter is the frequency domain representation of the steady-state KF.

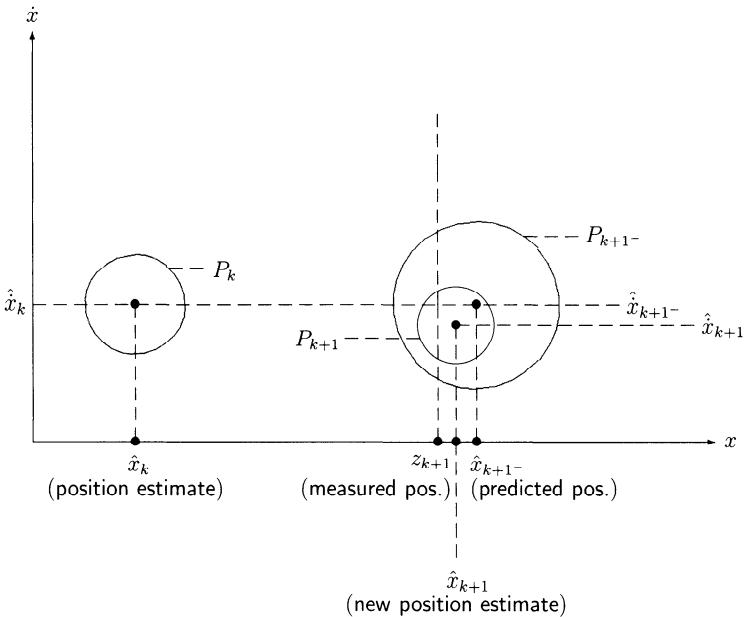


Figure 1.1.3-3: Illustration of the state estimation process.

The sampling interval (time between successive measurements) does not have to be constant in the KF.

Nonstationary disturbances (e.g., measurement noise with time-varying intensity) can be accommodated easily by the KF algorithm.

1.1.4 An Example of State Estimation: Vehicle Collision Avoidance

This subsection introduces the problem of state estimation for a dynamic system by considering the estimation of distance and speed between two vehicles as it arises in a collision avoidance system.

An algorithm is presented that estimates the distance (range) between two vehicles and their relative speed (range rate), based on imperfect (noisy) measurements of the distance between them. This algorithm is a simplified version of the Kalman filter, known as the $\alpha\text{-}\beta$ filter. It can also be viewed as a state observer [Chen84], which, based on observations of one state component, reconstructs (estimates) the entire state vector.

Define the state vector consisting of range r and range rate \dot{r} as

$$\boldsymbol{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} r \\ \dot{r} \end{bmatrix} \quad (1.1.4-1)$$

Assuming a constant speed (constant range-rate), the state equation is, for

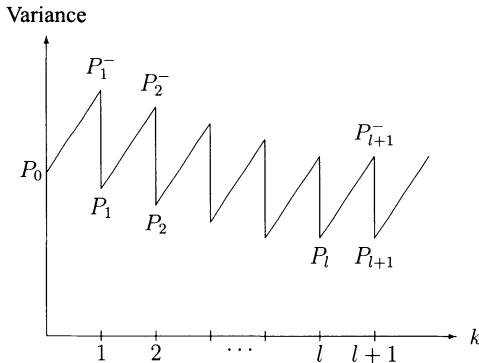


Figure 1.1.3-4: Evolution of the variance of the estimate with time.

discrete time with sampling period T ,

$$x(k) = \begin{bmatrix} r(k) \\ \dot{r}(k) \end{bmatrix} = \begin{bmatrix} 1 & T \\ 0 & 1 \end{bmatrix} x(k-1) = Fx(k-1) \quad (1.1.4-2)$$

where the system matrix is

$$F = \begin{bmatrix} 1 & T \\ 0 & 1 \end{bmatrix} \quad (1.1.4-3)$$

The measurement is

$$y(k) = x_1(k) = r(k) - [1 \ 0]x(k) = Hx(k) \quad (1.1.4-4)$$

where the measurement matrix is

$$H = [1 \ 0] \quad (1.1.4-5)$$

In truth, the measurement is a “corrupted” (noisy) version of (1.1.4-4), i.e.,

$$y(k) = Hx(k) + w(k) \quad (1.1.4-6)$$

where $w(k)$ is the (range) measurement error.

Similarly, the state equation should reflect an acceleration, denoted as u , namely,

$$x(k) = \begin{bmatrix} 1 & T \\ 0 & 1 \end{bmatrix} x(k-1) + \begin{bmatrix} T^2/2 \\ T \end{bmatrix} u(k-1) \quad (1.1.4-7)$$

Using (1.1.4-2) instead of (1.1.4-7), since we do not know u , as the “motion model” for the estimation problem leads to motion modeling errors.

The Estimator

In view of the fact that there are both measurement errors and motion modeling errors, we shall calculate the “estimate” $\hat{x}(k)$ of the state vector $x(k)$ via the

following algorithm, which is a simplified form of the Kalman filter, known as the α - β filter (see Chapters 5 and 6).

$$\hat{x}(k) = \bar{x}(k) + W[y(k) - \bar{y}(k)] \quad (1.1.4-8)$$

where $\bar{x}(k)$ is the “predicted state” at time k from $k-1$

$$\begin{aligned} \bar{x}(k) &= F\hat{x}(k-1) = \begin{bmatrix} 1 & T \\ 0 & 1 \end{bmatrix} \hat{x}(k-1) = \begin{bmatrix} 1 & T \\ 0 & 1 \end{bmatrix} \begin{bmatrix} \hat{r}(k-1) \\ \hat{\dot{r}}(k-1) \end{bmatrix} \\ &= \begin{bmatrix} \hat{r}(k-1) + T\hat{\dot{r}}(k-1) \\ \hat{\dot{r}}(k-1) \end{bmatrix} \end{aligned} \quad (1.1.4-9)$$

and $\bar{y}(k)$ is the “predicted measurement” at k from $k-1$

$$\bar{y}(k) = H\bar{x}(k) = [1 \ 0]\bar{x}(k) = \hat{r}(k-1) + T\hat{\dot{r}}(k-1) \quad (1.1.4-10)$$

The difference between the measurement (observed range) and its predicted value

$$y(k) - \bar{y}(k) = y(k) - H\bar{x}(k) \quad (1.1.4-11)$$

is the “residual,” which is multiplied (weighted) by the “filter gain”

$$W = \begin{bmatrix} \alpha \\ \beta/T \end{bmatrix} \quad (1.1.4-12)$$

The residual multiplied by the gain is used in (1.1.4-8) to “correct” or “update” the predicted state to obtain the new estimate $\hat{x}(k)$. The filter (1.1.4-8) can be seen as a ***predictor-corrector algorithm***.

The (dimensionless) constants α and β are such that $0 < \alpha < 1$ and $0 < \beta < 2$. The filter gain coefficients α and β define the α - β filter (1.1.4-8).

The breakdown of the state estimation equation (1.1.4-8) for its two components r and \dot{r} is

$$\hat{r}(k) = \hat{r}(k-1) + T\hat{\dot{r}}(k-1) + \alpha[y(k) - \bar{y}(k)] \quad (1.1.4-13)$$

$$\hat{\dot{r}}(k) = \hat{\dot{r}}(k-1) + \frac{\beta}{T}[y(k) - \bar{y}(k)] \quad (1.1.4-14)$$

The intuitive interpretation of (1.1.4-13) is as follows:

- If the predicted range is larger than the measured one, the residual will be negative and a negative correction is made with gain $\alpha < 1$ to obtain the new estimated range.
- If the predicted range is smaller than the measured one, the residual will be positive, and a positive correction will be made.

Similarly, the range rate (1.1.4-14) is corrected downward for a negative residual and upward for a positive one. Also note that the residual has the dimension of length and, when multiplied by β/T , its dimension becomes length/time, i.e., velocity.

A Deterministic Error Analysis

Assuming a constant acceleration (or deceleration) u , the state equation is, based on (1.1.4-7), given by

$$x(k) = Fx(k-1) + Gu \quad (1.1.4-15)$$

where

$$G = \begin{bmatrix} T^2/2 \\ T \end{bmatrix} \quad (1.1.4-16)$$

reflects the effect of the acceleration u on the range ($uT^2/2$) and on the range rate (uT).

Assuming the measurements to be given by (1.1.4-4), i.e., neglecting measurement errors, the evaluation of the effect of the acceleration on the errors in the range and range-rate estimates is carried out next.

The error is defined as the difference between the actual value and the estimate, namely,

$$\tilde{x}(k) = x(k) - \hat{x}(k) \quad (1.1.4-17)$$

From (1.1.4-8), using (1.1.4-9) and (1.1.4-10), the estimate is given by

$$\hat{x}(k) = F\hat{x}(k-1) + W[y(k) - HF\hat{x}(k-1)] \quad (1.1.4-18)$$

Substituting (1.1.4-4) for $y(k)$ in the above yields

$$\hat{x}(k) = F\hat{x}(k-1) + WH[x(k) - F\hat{x}(k-1)] \quad (1.1.4-19)$$

Substituting (1.1.4-15) for $x(k)$ in the above yields

$$\hat{x}(k) = F\hat{x}(k-1) + WH[Fx(k-1) + Gu - F\hat{x}(k-1)] \quad (1.1.4-20)$$

or, using notation (1.1.4-17),

$$\hat{x}(k) = F\hat{x}(k-1) + WHF\tilde{x}(k-1) + WHGu \quad (1.1.4-21)$$

Subtracting (1.1.4-21) from (1.1.4-15) yields

$$x(k) - \hat{x}(k) = Fx(k-1) + Gu - F\hat{x}(k-1) - WHF\tilde{x}(k-1) - WHGu \quad (1.1.4-22)$$

which can be rewritten as

$$\tilde{x}(k) = [I - WH]F\tilde{x}(k-1) + [I - WH]Gu \quad (1.1.4-23)$$

Note that the dynamic equation obeyed by the error \tilde{x} is a state equation with system matrix

$$\tilde{F} = [I - WH]F = \begin{bmatrix} 1 - \alpha & (1 - \alpha)T \\ -\frac{\beta}{T} & 1 - \beta \end{bmatrix} \quad (1.1.4-24)$$

and input with gain

$$\tilde{G} = [I - WH]G = \begin{bmatrix} (1-\alpha)T^2/2 \\ (1-\frac{\beta}{2})T \end{bmatrix} \quad (1.1.4-25)$$

i.e.,

$$\tilde{x}(k) = \tilde{F}\tilde{x}(k-1) + \tilde{G}u \quad (1.1.4-26)$$

It can be shown that the inequalities indicated earlier for α and β guarantee that the matrix \tilde{F} above is stable, i.e., it has eigenvalues inside the unit circle. In other words, for a constant “disturbance” u , the error $\tilde{x}(k)$ converges to a *steady-state* value \tilde{x}_{ss} , which is given by the expression

$$\tilde{x}_{ss} = (I - \tilde{F})^{-1}\tilde{G}u = \begin{bmatrix} \frac{1-\alpha}{\beta}T^2 \\ \frac{2\alpha-\beta}{2\beta}T \end{bmatrix} u \quad (1.1.4-27)$$

For $u = 5 \text{ m/s}^2$, $T = 0.1 \text{ s}$, $\alpha = 0.75$, $\beta = 0.5$, one obtains $\tilde{r}_{ss} = 0.39 \text{ m}$, $\dot{\tilde{r}}_{ss} = 0.5 \text{ m/s}$.

Frequency Domain View of Estimation: The Bandwidth of a Filter

Based on (1.1.4-23), one has the following z -domain transfer function from the input to the error \tilde{x} :

$$\tilde{H}(z) = \frac{\tilde{X}(z)}{U(z)} = (zI - \tilde{F})^{-1}\tilde{G} \quad (1.1.4-28)$$

With the above numbers,

$$\tilde{F} = \begin{bmatrix} 0.25 & 0.025 \\ -5 & 0.5 \end{bmatrix} \quad \tilde{G} = \begin{bmatrix} 0.005 \\ 0.1 \end{bmatrix} \quad (1.1.4-29)$$

the transfer function from the input to the position component error is

$$\tilde{H}_1(z) = \frac{\tilde{X}_1(z)}{U(z)} = \frac{0.005z}{z^2 - 0.75z + 0.25} \quad (1.1.4-30)$$

The two poles of $\tilde{H}_1(z)$ have magnitude $1/2$, i.e., this is a stable transfer function.

Figure 1.1.4-1 shows the plot of the magnitude of the above transfer function for the range of frequencies $[0, \frac{1}{2T}]$ (the upper limit is the point for which the Nyquist frequency is the sampling frequency). It can be clearly seen that this is a low-pass filter with a bandwidth (based on the half-power point) of approximately 2 Hz.

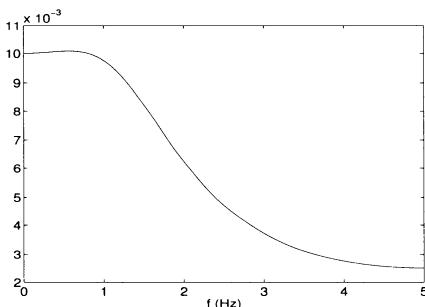


Figure 1.1.4-1: The frequency response of the state estimation filter.

1.2 SCOPE OF THE TEXT

1.2.1 Objectives

The objectives of this text are to present the fundamentals of *state estimation theory* and the tools for the design of state-of-the-art algorithms for *target tracking and navigation*.

The text covers the basic concepts and estimation techniques for static and dynamic systems, linear and nonlinear, as well as adaptive estimation. This constitutes a one-semester graduate course in estimation theory in an electrical/systems engineering program.

Special emphasis is given to the statistical tools that can be used for the *interpretation of the output from stochastic systems*. These are key tools for the assessment of the performance of state estimation/tracking filters — in simulations as well as in real-time implementation.

The discussion deals mainly with discrete time estimation algorithms, which are natural for digital computer implementation. Also, typically, sensors provide observations in discrete time. The basic state estimation algorithm — the Kalman filter — is presented in discrete as well as in continuous time.

The use of the estimation algorithms is illustrated on kinematic motion models because they reveal all the major issues and, in particular, the subtleties encountered in estimation. Also this serves as an introduction to tracking.

Guidelines for the *design of tracking filters*, namely,

- selection of the filter design parameters
- selection of the type of filter — nonadaptive (single model based) vs. adaptive (multiple model based)

are given and illustrated in several examples. Most of the examples are run with the companion software DynaEstTM; also, a step-by-step demonstration of how to use DynaEstTM is provided.

Prerequisite Material

The *linear algebra tools* that form the backbone of the state space analysis techniques used in estimation/tracking are presented in Section 1.3 together with some useful results from linear systems. Matrix notation will be used throughout this text — every quantity should be considered a matrix, unless otherwise stated.

A review of the *tools from probability theory and random processes*² is presented in Section 1.4.

Section 1.5 reviews *statistical hypothesis testing* and presents some useful statistical tables for the chi-square and the normal (Gaussian) distributions. Particular attention is given to the concept of statistical significance, which, while common in the statistics literature, has not received much attention in the engineering literature.

These review sections cover all the background material needed for the text and the basic concepts are illustrated by examples. This might be used as a (crash) minicourse to bring advanced undergraduate students up to the level necessary to understand the main material of the text.

1.2.2 Overview and Chapter Prerequisites

In the following an overview of the text is given with the prerequisite to each chapter.

Review of Background Techniques (remainder of Chapter 1)

- Linear algebra
- Probability theory and stochastic processes
- Statistics

Basic Concepts in Estimation — Chapter 2 (prerequisite: 1)

- Estimation techniques (ML, MAP, LS and MMSE)
- Properties of estimators
- Limit of existing information in the data

Linear Estimation in Static Systems — Chapter 3 (prerequisite: 2)

- Techniques (LMMSE, recursive LS and prelude to Kalman filtering)
- Application to polynomial fitting

Linear Dynamic Systems with Random Inputs — Chapter 4 (prerequisite: 1)

- Models — continuous and discrete time state space models
- Response of dynamic systems to random inputs

²Probability demonstrates both our knowledge as well as our ignorance.

State Estimation in Discrete-Time Linear Dynamic Systems — Chapter 5 (prerequisites: 3, 4)

- The Kalman filter and its properties
- Use of the companion software DynaEstTM to design and evaluate a KF
- Consistency of state estimators
- Initialization — in simulations and in practice

Estimation for Kinematic Models — Chapter 6 (prerequisite: 5)

- Types of kinematic models
- Explicit filters and the target maneuvering index
- Filter design and system performance prediction
- Sampling interval selection (system design)

Computational Aspects of Estimation — Chapter 7 (prerequisite: 5)

- The Information filter
- Sequential update
- Square-root filtering

Extensions of Discrete-Time Linear Estimation — Chapter 8 (prerequisite: 5)

- Correlated process noise
- Cross-correlated process and measurement noise sequences
- Correlated measurement noise
- Prediction
- Smoothing

Continuous-Time Linear State Estimation — Chapter 9 (prerequisite: 5)

- The Kalman-Bucy filter and its properties
- Prediction
- Duality between estimation and control
- The Wiener-Hopf problem

State Estimation for Nonlinear Dynamic Systems — Chapter 10 (prerequisite: 5)

- Optimal estimation
- Suboptimal estimator — the extended Kalman filter (EKF)
- Practical issues in implementation
- Trajectory estimation via dynamic programming

Adaptive Estimation and Maneuvering Targets — Chapter 11 (prerequisite: 5)

- White and colored noise models
- Input estimation
- Variable state dimension filtering

- Comparison of several methods
- The multiple model approach — static and dynamic, the Interacting Multiple Model (IMM) estimator
- Use of the companion software DynaEstTM to design and evaluate an IMM estimator
- An air traffic control (ATC) tracking algorithm design example with the IMM estimator
- Estimator selection — nonadaptive vs. adaptive
- Use of the EKF for parameter estimation

Estimation for Navigation Systems — Chapter 12 (prerequisites: 5 and, for the last part, 11)

- Inertial navigation systems (INS)
- Global Positioning System (GPS)
- Combination of INS and GPS
- Use of the KF and IMM estimator for GPS navigation

At the end of each chapter, a number of problems that enhance the understanding of the theory and connect the theoretical material to the real world are given.

The sequence of prerequisites is illustrated in Fig. 1.2.2-1.

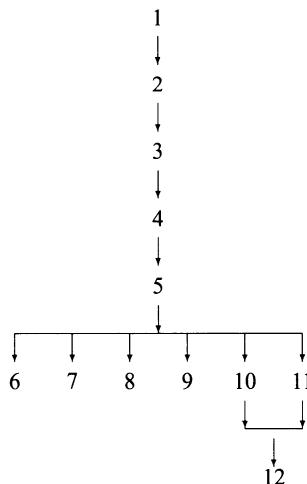


Figure 1.2.2-1: Sequence of chapter prerequisites.

Two major examples of realistic systems are presented:

- A bearings-only target localization from a moving platform based on the Maximum Likelihood technique
- An ATC problem for tracking maneuvering aircraft using an adaptive estimation technique (the IMM estimator)

As part of the problems at the end of Chapter 3, a project is given to develop an interactive program for the first example listed above.

A major term project that involves a realistic ATC tracking problem (with nonlinear measurements from a radar — range and azimuth) is given at the end of Chapter 11. The sequence of prerequisites for this is Chapters 1–6, Section 10.3 and Section 11.6.

An alternative major term project that involves a realistic ATC navigation problem with GPS is given at the end of Chapter 11. The sequence of prerequisites for this is Chapters 1–6, Section 10.3 and Sections 12.6–12.8.

1.3 BRIEF REVIEW OF LINEAR ALGEBRA AND LINEAR SYSTEMS

1.3.1 Definitions and Notations

A **matrix** of dimension $n \times m$ is the two-dimensional array

$$A = [a_{ij}] = \begin{bmatrix} a_{11} & \cdots & a_{1m} \\ \vdots & \ddots & \vdots \\ a_{n1} & \cdots & a_{nm} \end{bmatrix} \quad (1.3.1-1)$$

The first dimension is the number of rows, and the second dimension is the number of columns. The elements of A will be denoted as a_{ij} or as A_{ij} .

An n -dimensional **vector** is the one-dimensional array ($n \times 1$ matrix)

$$a = \begin{bmatrix} a_1 \\ \vdots \\ a_n \end{bmatrix} \triangleq \text{col}(a_i) \quad (1.3.1-2)$$

In this text all vectors will be **column vectors**.

Transposition of a matrix or vector will be denoted by an apostrophe. The **transpose** of the matrix (1.3.1-1) is

$$A' = [a_{ji}] = \begin{bmatrix} a_{11} & \cdots & a_{n1} \\ \vdots & \ddots & \vdots \\ a_{1m} & \cdots & a_{nm} \end{bmatrix} \quad (1.3.1-3)$$

With this notation one can write the column vector (1.3.1-2) as

$$a = [a_1 \ \cdots \ a_n]' \quad (1.3.1-4)$$

The **row vector** obtained from transposing (1.3.1-2) is

$$a' = [a_1 \ \cdots \ a_n] \quad (1.3.1-5)$$

A (square) matrix is said to be **symmetric** if

$$A = A' \quad (1.3.1-6)$$

or

$$a_{ij} = a_{ji} \quad \forall i, j \quad (1.3.1-7)$$

1.3.2 Some Linear Algebra Operations

Addition of matrices and the multiplication of a matrix by a scalar are defined as follows. With α and β scalars, the matrix

$$C = \alpha A + \beta B \quad (1.3.2-1)$$

has elements given by

$$c_{ij} = \alpha a_{ij} + \beta b_{ij} \quad i = 1, \dots, n; \quad j = 1, \dots, m \quad (1.3.2-2)$$

where all the matrices have the same dimension $n \times m$.

The product of two matrices

$$C = AB \quad (1.3.2-3)$$

has elements

$$c_{ij} = \sum_{k=1}^m a_{ik} b_{kj} \quad i = 1, \dots, n; \quad j = 1, \dots, p \quad (1.3.2-4)$$

where A is $n \times m$, B is $m \times p$, and the result C is $n \times p$. The matrix product is, in general, not commutative.

The **transpose of a product** is

$$C' = (AB)' = B'A' \quad (1.3.2-5)$$

Thus, if

$$Ab = c \quad (1.3.2-6)$$

where A is $n \times m$, b is $m \times 1$ (i.e., an m -vector) and c is $n \times 1$ (n -vector), then

$$c' = b'A' \quad (1.3.2-7)$$

where c' is $1 \times n$ (n -vector in row form), etc. The **inner product** of two (real) n -vectors in a Euclidean space is

$$\langle a, b \rangle \triangleq a'b = \sum_{i=1}^n a_i b_i \quad (1.3.2-8)$$

The inner product has a counterpart in probability theory where it is used for random variables, which can be looked upon as vectors in a suitable space. This has applications in linear minimum mean square error estimation.

The **outer product** of two vectors is the matrix

$$ab' = C = [c_{ij}] \quad \text{with} \quad c_{ij} = a_i b_j \quad (1.3.2-9)$$

The **trace** of an $n \times n$ matrix is the sum of its diagonal elements

$$\text{tr}(A) \triangleq \sum_{i=1}^n a_{ii} = \text{tr}(A') \quad (1.3.2-10)$$

It can be easily shown³ that if A is $m \times n$ and B is $n \times m$, then

$$\text{tr}(AB) = \text{tr}(BA) \quad (1.3.2-11)$$

For example, if a and b are both n -vectors, then

$$a'b = \text{tr}(a'b) = \text{tr}(ba') \quad (1.3.2-12)$$

The first equality above is due to the fact that the trace of a scalar is itself.

It can be shown⁴ that for matrices of suitable dimension,

$$\text{tr}(ABC) = \text{tr}(BCA) = \text{tr}(CAB) \quad (1.3.2-13)$$

that is, the trace is *invariant under circular permutations* in its argument. A (square) matrix is ***idempotent*** if and only if⁵ for all positive integers n

$$A^n = A \quad (1.3.2-14)$$

The ***square root*** of the square matrix A is the (in general, nonunique) matrix $A^{1/2}$, such that

$$A^{1/2}A^{1/2} = A \quad (1.3.2-15)$$

A symmetric matrix A can also be decomposed as

$$A = CC' = LDL' \quad (1.3.2-16)$$

where C is ***lower triangular*** (with zeros above the diagonal), L is ***unit lower triangular*** (with zeros above the diagonal and units on the diagonal) and D is ***diagonal***. The first decomposition above is the ***Cholesky factorization*** whereas the last form above is known as the ***LDL'* factorization**, or, since $L' = U$ is a unit upper triangular matrix, the ***LDU factorization***.

1.3.3 Inversion and the Determinant of a Matrix

The ***inverse*** A^{-1} of the $n \times n$ matrix A is such that

$$A^{-1}A = I \quad (1.3.3-1)$$

where

$$I = \begin{bmatrix} 1 & 0 & \cdots & 0 \\ 0 & 1 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & 1 \end{bmatrix} \triangleq \text{diag}(1) \quad (1.3.3-2)$$

³Also known in mathematical circles as the ICBES argument.

⁴Also known in mathematical circles as the ICBS argument.

⁵The standard shorthand for “if and only if” in mathematical texts is “***iff***.”

is the **identity matrix** of the same dimension as A .

The inverse is given by the expression

$$A^{-1} = \frac{1}{|A|} C' \quad (1.3.3-3)$$

where C , called the **adjugate** of A , is the matrix of cofactors of A and $|A|$ is the **determinant** of A . The cofactor (i, j) of A is the determinant of the matrix A without its row i and column j , multiplied by $(-1)^{i+j}$.

The inverse of a (square) matrix exists if and only if its columns a_i , $i = 1, \dots, n$ (or its rows) are **linearly independent**, that is,

$$\sum_{i=1}^n \alpha_i a_i = 0 \iff \alpha_i = 0 \quad i = 1, \dots, n \quad (1.3.3-4)$$

where 0 denotes the zero vector (of dimension n). This is equivalent to the determinant in (1.3.3-3) being nonzero.

An invertible matrix is also called **nonsingular**.

The determinant of an $n \times n$ matrix multiplied by a scalar is

$$|\alpha A| = \alpha^n |A| \quad (1.3.3-5)$$

This is useful in writing the probability density function of Gaussian random vector.

Inversion of a Partitioned Matrix

The inverse of the (nonsingular) $n \times n$ **partitioned matrix**

$$\begin{bmatrix} P_{11} & P_{12} \\ P_{21} & P_{22} \end{bmatrix}^{-1} = \begin{bmatrix} V_{11} & V_{12} \\ V_{21} & V_{22} \end{bmatrix} \quad (1.3.3-6)$$

where P_{11} is $n_1 \times n_1$, P_{12} is $n_1 \times n_2$, P_{21} is $n_2 \times n_1$, P_{22} is $n_2 \times n_2$ and $n_1 + n_2 = n$, has the partitions

$$V_{11} = P_{11}^{-1} + P_{11}^{-1} P_{12} V_{22} P_{21} P_{11}^{-1} = (P_{11} - P_{12} P_{22}^{-1} P_{21})^{-1} \quad (1.3.3-7)$$

$$V_{12} = -P_{11}^{-1} P_{12} V_{22} = -V_{11} P_{12} P_{22}^{-1} \quad (1.3.3-8)$$

$$V_{21} = -V_{22} P_{21} P_{11}^{-1} = -P_{22}^{-1} P_{21} V_{11} \quad (1.3.3-9)$$

$$V_{22} = P_{22}^{-1} + P_{22}^{-1} P_{21} V_{11} P_{12} P_{22}^{-1} = (P_{22} - P_{21} P_{11}^{-1})^{-1} P_{12} \quad (1.3.3-10)$$

The Matrix Inversion Lemma

Another useful result is the following identity known as the ***matrix inversion lemma***,

$$(P^{-1} + H'R^{-1}H)^{-1} = P - PH'(HPH' + R)^{-1}HP \quad (1.3.3-11)$$

where P is $n \times n$, H is $m \times n$, and R is $m \times m$.

An alternative version of the above is

$$(A + BCB')^{-1} = A^{-1} - A^{-1}B(B'A^{-1}B + C^{-1})^{-1}B'A^{-1} \quad (1.3.3-12)$$

It can be shown that (1.3.3-6) to (1.3.3-10) and (1.3.3-11) hold by verifying that the corresponding multiplications will yield the identity matrix.

These results have direct application in the derivation of the recursive form of the least squares estimation of parameters as well as in linear estimation for dynamic systems — they yield various forms of the Kalman filter.

1.3.4 Orthogonal Projection of Vectors

The inner product of a vector with itself

$$\langle a, a \rangle \triangleq \|a\|^2 \quad (1.3.4-1)$$

is the **squared l_2 norm** of this vector. This applies for the inner product defined in (1.3.2-8) or for any other properly defined inner product. The l_2 norm is the **magnitude** of the vector or its **length**.

The **Schwarz inequality** states the following relationship between the magnitude of the inner product of two vectors and their norms

$$|\langle a, b \rangle| \leq \|a\| \|b\| \quad (1.3.4-2)$$

Two vectors are **orthogonal**, which is denoted as

$$a \perp b \quad (1.3.4-3)$$

if

$$\langle a, b \rangle = 0 \quad (1.3.4-4)$$

The **orthogonal projection** of the vector a on b is

$$\Pi_b(a) = \frac{\langle a, b \rangle}{\|b\|^2} b \quad (1.3.4-5)$$

It can be easily shown that the difference vector between a and its orthogonal projection on b is orthogonal on b , that is,

$$[a - \Pi_b(a)] \perp b \quad (1.3.4-6)$$

Figure 1.3.4-1 illustrates this.

The orthogonal projection is used to obtain the linear minimum mean square error (LMMSE) estimates of random variables.

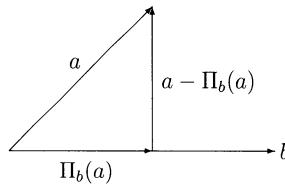


Figure 1.3.4-1: Orthogonal projection of vector a on vector b .

1.3.5 The Gradient, Jacobian and Hessian

The **gradient** operator with respect to the n -vector x is defined as

$$\nabla_x = \left[\frac{\partial}{\partial x_1} \cdots \frac{\partial}{\partial x_n} \right]' \quad (1.3.5-1)$$

The gradient of an m -dimensional vector-valued function $f(x)$ is

$$\nabla_x f(x)' = \begin{bmatrix} \frac{\partial}{\partial x_1} \\ \vdots \\ \frac{\partial}{\partial x_n} \end{bmatrix} [f_1(x) \cdots f_m(x)] = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \cdots & \frac{\partial f_m}{\partial x_1} \\ \vdots & \ddots & \vdots \\ \frac{\partial f_1}{\partial x_n} & \cdots & \frac{\partial f_m}{\partial x_n} \end{bmatrix} \quad (1.3.5-2)$$

The transpose of the above is the **Jacobian matrix**, an $m \times n$ matrix in this case

$$f_x(x) \triangleq \frac{\partial f}{\partial x} = [\nabla_x f(x)']' = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \cdots & \frac{\partial f_1}{\partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial f_m}{\partial x_1} & \cdots & \frac{\partial f_m}{\partial x_n} \end{bmatrix} \quad (1.3.5-3)$$

The dimensions m and n are usually (but not necessarily) equal.

It can be easily shown that

$$\nabla_x x' = I \quad (1.3.5-4)$$

and, if A is a symmetric matrix,

$$\nabla_x (x' A x) = 2Ax \quad (1.3.5-5)$$

The **Hessian** of the scalar function $\phi(x)$ with respect to the n -vector x is

$$\phi_{xx}(x) \triangleq \frac{\partial^2 \phi(x)}{\partial x^2} = \nabla_x \nabla'_x \phi(x) = \begin{bmatrix} \frac{\partial^2 \phi}{\partial x_1 \partial x_1} & \cdots & \frac{\partial^2 \phi}{\partial x_1 \partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial^2 \phi}{\partial x_n \partial x_1} & \cdots & \frac{\partial^2 \phi}{\partial x_n \partial x_n} \end{bmatrix} \quad (1.3.5-6)$$

which is, obviously, a symmetric $n \times n$ matrix.

These results are used in the linear estimation of the state of nonlinear systems via series expansion.

1.3.6 Eigenvalues, Eigenvectors, and Quadratic Forms

Eigenvalues and Eigenvectors

The **eigenvalues** of an $n \times n$ matrix A are the scalars λ_i such that

$$Au_i = \lambda_i u_i \quad i = 1, \dots, n \quad (1.3.6-1)$$

where the vectors u_i are the corresponding **eigenvectors**. If the eigenvalues are distinct, then there are n eigenvectors; otherwise the number of the linearly independent eigenvectors might be less than n .

A matrix is **nonsingular** if and only if all its eigenvalues are nonzero. This follows from the fact that its determinant is the product of its eigenvalues

$$|A| = \prod_{i=1}^n \lambda_i \quad (1.3.6-2)$$

The **rank** of a (not necessarily square) matrix is equal to the number of its linearly independent rows or columns. For a square ($n \times n$) matrix, this equals the number of its nonzero eigenvalues.

A nonsingular matrix must have **full rank**. In the case of a full rank $n \times n$ matrix, its columns (and rows) span the n -dimensional space.

The eigenvalues of a real matrix can be real or complex; however, a *symmetric* real matrix has only real eigenvalues.

It can be shown that the trace of a matrix equals the sum of its eigenvalues

$$\text{tr}(A) = \sum_{i=1}^n \lambda_i \quad (1.3.6-3)$$

An **idempotent** matrix, defined in (1.3.2-14), that is, with the property

$$AA = A \quad (1.3.6-4)$$

has eigenvalues that are *either zero or unity*.

Quadratic Forms

The (scalar) function of the real vector x

$$q = x'Ax \quad (1.3.6-5)$$

is called a **quadratic form**. It can be easily shown that, without loss of generality, in a quadratic form the matrix A can be considered symmetric.

A **positive (semi)definite quadratic form** is one which is positive (nonnegative) for all nonzero vectors x .

A **positive (semi)definite matrix** is one for which the quadratic form (1.3.6-5) is positive (semi)definite. This can be summarized as follows:

$$A > 0 \iff x'Ax > 0 \quad \forall x \neq 0 \quad (1.3.6-6)$$

$$A \geq 0 \iff x'Ax \geq 0 \quad \forall x \neq 0 \quad (1.3.6-7)$$

If the matrix A is positive definite, the expression on the right hand side of (1.3.6-5) can also be called the **squared norm with respect to A** of the vector x and denoted sometimes as $\|x\|_A^2$.

A matrix is positive (semi)definite if and only if all its eigenvalues are positive (nonnegative).

The **inequality of two matrices** is defined as follows: the matrix A is smaller (not larger) than the matrix B if and only if the difference $B - A$ is positive (semi)definite. The outer product of a vector with itself, aa' , called a **dyad**, is a positive semidefinite matrix (it has rank one).

The sum of dyads of n -vectors a_i with positive weights

$$\tilde{A} = \sum_{i=1}^m \alpha_i a_i a_i' \quad \alpha_i > 0 \quad (1.3.6-8)$$

is (at least) positive semidefinite. If the vectors a_1, \dots, a_m span the n -dimensional space, in which case one needs $m \geq n$, then (1.3.6-8) is positive definite. The **spectral representation** of an $n \times n$ matrix that has n eigenvectors is

$$A = \sum_{i=1}^n \lambda_i u_i v_i' \quad (1.3.6-9)$$

where λ_i are its eigenvalues, u_i are its eigenvectors, and v_i , $i = 1, \dots, n$, is the **reciprocal basis**, that is,

$$u_i' v_j = \delta_{ij} = \begin{cases} 1 & \text{if } i = j \\ 0 & \text{if } i \neq j \end{cases} \quad (1.3.6-10)$$

The vectors v_i are the columns of the inverse of the matrix consisting of the eigenvectors u_i' as rows.

If the $n \times n$ real matrix A is **symmetric** with nonzero eigenvalues, it will always have n eigenvectors that are **orthogonal**. Furthermore, in this case,

$$v_i = u_i \quad (1.3.6-11)$$

Thus, the spectral representation of a **symmetric** real $n \times n$ matrix A is

$$A = \sum_{i=1}^n \lambda_i u_i u_i' \quad (1.3.6-12)$$

The **condition number** of a positive definite symmetric matrix is (usually taken as) the common logarithm of the ratio of its largest to its smallest eigenvalue:

$$\kappa(A) \triangleq \log_{10} \frac{\lambda_{\max}}{\lambda_{\min}} \quad (1.3.6-13)$$

Large condition numbers indicate near-singularity (e.g., $\kappa > 6$ for a 32-bit computer indicates an ill-conditioned matrix).

The (induced) **norm of a matrix** is

$$\|A\| \triangleq \max_{x \neq 0} \frac{\|Ax\|}{\|x\|} \quad (1.3.6-14)$$

A consequence of the above and the Schwarz inequality is

$$\|Ax\| \leq \|A\| \|x\| \quad (1.3.6-15)$$

Also

$$\|AB\| \leq \|A\| \|B\| \quad (1.3.6-16)$$

1.3.7 Continuous-Time Linear Dynamic Systems — Controllability and Observability

The State-Space Representation

The state space representation of *continuous-time linear systems* is

$$\dot{x}(t) = A(t)x(t) + B(t)u(t) \quad (1.3.7-1)$$

where

$x(t)$ is the state vector of dimension n_x ,

$u(t)$ is the input (control) vector of dimension n_u ,

$A(t)$ is an $n_x \times n_x$ matrix, called the **system matrix**,

$B(t)$ is an $n_x \times n_u$ matrix, called the (continuous-time) **input gain**.

Equation (1.3.7-1) is known as the **dynamic equation** or the **plant equation**.

The output of the system is, in general, a vector

$$z(t) = C(t)x(t) \quad (1.3.7-2)$$

of dimension n_z and C a known $n_z \times n_x$ matrix, called the **measurement matrix**.

Equation (1.3.7-2) is known as the **output equation** or the **measurement equation**.

Given the initial condition $x(t_0)$ and the input function denoted as

$$u_{[t_0,t]} = \{u(\tau), t_0 \leq \tau \leq t\} \quad (1.3.7-3)$$

one can compute the future output at any time $t > t_0$

$$z(t) = z[x(t_0), u_{[t_0,t]}, t, t_0] \quad (1.3.7-4)$$

The **state** (of a deterministic system) is defined as the smallest dimension vector that *completely summarizes the past of the system*.

A continuous-time system is **linear time-invariant** if it is described by (1.3.7-1) and (1.3.7-2) with $A(t) = A$, $B(t) = B$, and $C(t) = C$.

Controllability

A continuous-time (deterministic) system is **completely controllable** if, given an *arbitrary destination point* in the state space, there is an input function that will bring the system from any initial state to this point in a finite time.

For a *linear time-invariant system*, the controllability condition is that the pair $\{A, B\}$ is controllable, that is, the **controllability matrix**

$$\mathcal{Q}_C \triangleq [B \ AB \ \cdots \ A^{n_x-1}B] \quad (1.3.7-5)$$

has **full rank**, which in this case is n_x , the lower of its two dimensions (the other is $n_x n_u$).

Observability

A continuous-time (deterministic) system is **completely observable** if its initial state can be *fully and uniquely* recovered from its output, observed over a finite time interval, and the knowledge of the input.

Note that since the system is deterministic, knowledge of the initial state is equivalent to knowledge of the state at any time. Thus, using the output (1.3.7-2), which is in general a vector of dimension $n_z < n_x$, and the input, one can then recover the state perfectly.

For a *linear time-invariant system*, the observability condition is that the pair $\{A, C\}$ is observable; that is, the **observability matrix**

$$\mathcal{Q}_O \triangleq \begin{bmatrix} C \\ CA \\ \vdots \\ CA^{n_x-1} \end{bmatrix} \quad (1.3.7-6)$$

has full rank n_x .

Example

Consider the system (double integrator)

$$\ddot{\xi}(t) = u(t) \quad (1.3.7-7)$$

With the state vector consisting of the position ξ and velocity $\dot{\xi}$

$$x = [\xi \ \dot{\xi}]' \quad (1.3.7-8)$$

the state space representation of this system is

$$\dot{x}(t) = Ax(t) + Bu(t) = \begin{bmatrix} 0 & 1 \\ 0 & 0 \end{bmatrix} x(t) + \begin{bmatrix} 0 \\ 1 \end{bmatrix} u(t) \quad (1.3.7-9)$$

The controllability matrix is

$$\mathcal{Q}_C = [B \ AB] = \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix} \quad (1.3.7-10)$$

which has full rank $n_x = 2$, i.e., the state is *completely controllable*. While the input u enters directly only into the velocity component of the state, it does have the capability of controlling both state components. If the observation is

$$z(t) = Cx(t) = [1 \ 0]x(t) \quad (1.3.7-11)$$

i.e., the position ξ , the observability matrix is

$$\mathcal{Q}_O = \begin{bmatrix} C \\ CA \end{bmatrix} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \quad (1.3.7-12)$$

and it has full rank $n_x = 2$, i.e., the state is *completely observable*.

If the observation is

$$z(t) = Cx(t) = [0 \ 1]x(t) \quad (1.3.7-13)$$

i.e., the velocity $\dot{\xi}$, the observability matrix is

$$\mathcal{Q}_O = \begin{bmatrix} C \\ CA \end{bmatrix} = \begin{bmatrix} 0 & 1 \\ 0 & 0 \end{bmatrix} \quad (1.3.7-14)$$

and it has rank $1 < n_x = 2$, i.e., the state is *not completely observable*. The reason is that, while the position observations (1.3.7-11) provide information about the velocity, the velocity observations (1.3.7-13) do not provide any position information.

1.3.8 Discrete-Time Linear Dynamic Systems — Controllability and Observability

The State-Space Representation

Deterministic linear dynamic systems can be described in discrete time by a state equation of the form

$$x(k+1) = F(k)x(k) + G(k)u(k) \quad (1.3.8-1)$$

where

- $x(k)$ is the state of the system, a vector of dimension n_x ,
- $F(k)$ is the transition matrix ($n_x \times n_x$),
- $u(k)$ is the input (control), a vector of dimension n_u ,
- $G(k)$ is the input gain ($n_x \times n_u$ matrix),

all at time k .

The output equation is

$$z(k) = H(k)x(k) \quad (1.3.8-2)$$

where

$z(k)$ is the output vector (observation or measurement), of dimension n_z ,
 $H(k)$ is the measurement matrix ($n_z \times n_x$).

The **state** of the system is the smallest dimension vector that completely summarizes the past of the system. Then the output at time $j > k$ can be determined fully from $x(k)$ and the intervening inputs

$$z(j) = z[j, x(k), u(k), \dots, u(j-1)] \quad (1.3.8-3)$$

A discrete-time system is **linear time-invariant** if it is described by (1.3.8-1) and (1.3.8-2) with $F(k) = F$, $G(k) = G$, and $H(k) = H$.

Controllability

A discrete-time (deterministic) system is **completely controllable** if, given an *arbitrary destination point* in the state space, there is an input sequence that will bring the system from any initial state to this point in a finite number of steps.

For a *linear time-invariant system*, the controllability condition is that the pair $\{F, G\}$ is controllable; that is, the **controllability matrix**

$$\mathcal{Q}_C \triangleq [G \ FG \ \cdots \ F^{n_x-1}G] \quad (1.3.8-4)$$

has full rank n_x .

Observability

A (deterministic) system is **completely observable** if its initial state can be *fully and uniquely* recovered from a finite number of observations of its output and the knowledge of its input.

Note that since the system is deterministic, knowledge of the initial state is equivalent to knowledge of the state at any time. Thus, using the input and the output (1.3.8-2), which is in general a vector of dimension $n_z < n_x$, one can, *in the absence of noise*, recover the state perfectly.

For a *linear time-invariant system*, the observability condition is that the pair $\{F, H\}$ is observable; that is, the **observability matrix**

$$\mathcal{Q}_O \triangleq \begin{bmatrix} H \\ HF \\ \vdots \\ HF^{n_x-1} \end{bmatrix} \quad (1.3.8-5)$$

has full rank n_x .

1.4 BRIEF REVIEW OF PROBABILITY THEORY

1.4.1 Events and the Axioms of Probability

Consider an “experiment,” or, in general, a process with random outcomes. An **event** is a collection (set) of such outcomes — it is said to have occurred if the outcome is one of the elements of this set.

Denote by A an event in such an experiment (e.g., “even” or “5” in a die rolling experiment). Let S be the **sure event** in the experiment (e.g., “any number between 1 and 6” in the die rolling). Then the **probability** of an event is a number (**measure**) that satisfies the following three **axioms of probability**:

1. It is nonnegative

$$P\{A\} \geq 0 \quad \forall A \quad (1.4.1-1)$$

2. It is unity for the sure event

$$P\{S\} = 1 \quad (1.4.1-2)$$

3. It is additive over the union of **mutually exclusive** events; that is, if the events A and B have no common elements (their set intersection is \emptyset , the empty set)

$$A \cap B \triangleq \{A \text{ and } B\} \triangleq \{A, B\} = \emptyset \quad (1.4.1-3)$$

then their union (logical “or”) has probability

$$P\{A \cup B\} \triangleq P\{A \text{ or } B\} \triangleq P\{A + B\} = P\{A\} + P\{B\} \quad (1.4.1-4)$$

From the above it follows that

$$P\{A\} = 1 - P\{\bar{A}\} \leq 1 \quad (1.4.1-5)$$

where the overbar denotes the **complementary event**.

Since \emptyset is the complement of S , one has

$$P\{\emptyset\} = 0 \quad (1.4.1-6)$$

The event \emptyset is called the **impossible event**.

Extension

The extended version of Axiom 3 such that the probability is additive over the union of an *infinite number of mutually exclusive events* — **countable additivity** — is necessary when an experiment has an infinite number of outcomes. This is a key point in the definition of continuous-valued random variables.

Remarks

Relative-frequency or **measure-of-belief** interpretation of probability are alternative, even though not rigorous, ways of introducing the concept of probability.

In practice one can use the relative frequency or measure of belief to assign, based on intuition (or engineering common sense), probabilities to certain events. This can be of major importance in certain engineering systems.

Note that probability is a scalar quantity that is dimensionless (in a physical sense). This is in contrast to, for instance, a probability density function (pdf) which, while a scalar, has in general a *physical dimension*.

Example

For a die-rolling experiment, the space of outcomes is

$$S = \{f_1, \dots, f_6\} \quad (1.4.1-7)$$

where f_i denotes “face i ”, $i = 1, \dots, 6$.

Events are subsets of S , for example,

$$F_1 = \{f_1\} \quad (1.4.1-8)$$

$$\{i = \text{even}\} = \{f_2, f_4, f_6\} \quad (1.4.1-9)$$

$$\{2 < i \leq 4\} = \{f_3, f_4\} \quad (1.4.1-10)$$

The sure event

$$\{\text{any face}\} = S \quad (1.4.1-11)$$

is denoted by the same symbol as the space of outcomes since it is the set of all possible outcomes.

An impossible event is, for example,

$$\{i \leq 1 \text{ and even}\} = \emptyset \quad (1.4.1-12)$$

This is clearly an empty subset of S .

Probabilities are assigned (postulated) for each element of S (each outcome). For a finite number of outcomes this *granular approach* yields, using Axiom 3, the probabilities of all the subsets of S , i.e., the events.

Postulating a fair die — each outcome (face) is equally probable — one has from Axioms 2 and 3,

$$1 = P\{S\} = P\left\{\bigcup_{i=1}^6 F_i\right\} = \sum_{i=1}^6 P\{F_i\} = 6P\{F_i\} \quad (1.4.1-13)$$

which yields

$$P\{F_i\} = 1/6 \quad i = 1, \dots, 6 \quad (1.4.1-14)$$

From this it follows with Axiom 3 that

$$P\{i = \text{even}\} = P\{F_2\} + P\{F_4\} + P\{F_6\} = 1/2 \quad (1.4.1-15)$$

1.4.2 Random Variables and Probability Density Function

A scalar **random variable** is a (real-valued) function that assumes a certain *value* according to the outcome of a *random experiment*. The value taken by a random variable is called its **realization**.

The **probability density function (pdf)** of a scalar continuous-valued random variable x at $x = \xi$ is

$$p_x(\xi) = \lim_{d\xi \rightarrow 0} \frac{P\{\xi - d\xi < x \leq \xi\}}{d\xi} \geq 0 \quad (1.4.2-1)$$

where $P\{\cdot\}$ is the probability of the event $\{\cdot\}$.

The more common notation

$$p_x(x) = p(x) \quad (1.4.2-2)$$

where the argument defines the function, is used instead of (1.4.2-1). Also the term **density** is sometimes used instead of *pdf*.

From (1.4.2-1) and Axiom 3 from Subsection 1.4.1 it follows that

$$P\{\eta < x \leq \xi\} = \int_{\eta}^{\xi} p(x) dx \quad (1.4.2-3)$$

The function

$$P_x(\xi) = P\{x \leq \xi\} = \int_{-\infty}^{\xi} p(x) dx \quad (1.4.2-4)$$

is called the **cumulative probability distribution function (cdf)** of x at ξ . This is usually referred to as **distribution**.

Since the event $\{x \leq \infty\}$ is the “sure” event, one has

$$P\{x \leq \infty\} = \int_{-\infty}^{\infty} p(x) dx = 1 \quad (1.4.2-5)$$

A pdf has to have the **normalization property** (1.4.2-5) that its total **probability mass** is unity — otherwise it is not a **proper density**.

The relationship between the density and the cumulative distribution is, from (1.4.2-4),

$$p(x) = \left. \frac{d}{d\xi} P_x(\xi) \right|_{\xi=x} \quad (1.4.2-6)$$

if the derivative exists.

An outline of the rigorous way of introducing the pdf is as follows:

1. First define the “basic” events $\{x \leq \xi\}$.
2. Show that from these events one can obtain via (a countable number of) set operations all the events of interest (e.g., $\{a < x \leq \xi\}$, $\{x = a\}$).
3. Define the probabilities of the basic events, from which the probabilities of the other events of interest can be computed using the extended version of the third axiom of probability.
4. Finally, the pdf follows from (1.4.2-6) under suitable conditions of differentiability; note that (1.4.2-6) is equivalent to (1.4.2-1).

Remarks

Note from (1.4.2-1) that with the numerator being dimensionless, the pdf of the random variable x has as its *physical dimension* the inverse of the physical dimension of x . (See also problem 1-5.)

The event $\{x = a\}$ has, for a continuous-valued random variable, probability zero (even though it is *not impossible!*). This follows from (1.4.2-3) when the interval length tends to zero, in which case, as long as the density is finite, the integral is zero. Finite density means “no point masses” — this is discussed in the next section.

Improper densities — which do not integrate to unity (actually, their integral is not defined) — can be used, however, in certain circumstances.

Example

A **uniformly distributed random variable** in the interval $[a, b]$, symbolically denoted as

$$x \sim \mathcal{U}(a, b) \quad (1.4.2-7)$$

has the pdf

$$p(x) = \mathcal{U}(x; a, b) \triangleq \begin{cases} \frac{1}{b-a} & x \in [a, b] \\ 0 & \text{elsewhere} \end{cases} \quad (1.4.2-8)$$

Its cdf is

$$P_x(\xi) = \begin{cases} 0 & \xi < a \\ \frac{\xi - a}{b - a} & \xi \in [a, b] \\ 1 & \xi > b \end{cases} \quad (1.4.2-9)$$

A **Gaussian (normal) random variable** with parameters μ, σ^2 (the mean and variance, respectively, discussed later) is symbolically denoted as

$$x \sim \mathcal{N}(\mu, \sigma^2) \quad (1.4.2-10)$$

Its pdf is

$$p(x) = \mathcal{N}(x; \mu, \sigma^2) \triangleq \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}} \quad (1.4.2-11)$$

and the cdf is

$$\begin{aligned} P_x(\xi) &= \int_{-\infty}^{\xi} \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}} dx = \int_{-\infty}^{\frac{\xi-\mu}{\sigma}} \frac{1}{\sqrt{2\pi}} e^{-\frac{\eta^2}{2}} d\eta \\ &= \int_{-\infty}^{\frac{\xi-\mu}{\sigma}} \mathcal{N}(\eta; 0, 1) d\eta \triangleq \mathcal{G}\left(\frac{\xi-\mu}{\sigma}\right) \end{aligned} \quad (1.4.2-12)$$

where \mathcal{G} is the **cumulative standard Gaussian distribution** (with mean $\mu = 0$ and standard deviation $\sigma = 1$). Its argument above is the *number of standard deviations away from the mean*. This integral has been extensively tabulated (see, e.g., Table 1.5.4-1).

Note

The symbols \mathcal{U} and \mathcal{N} with two arguments indicate “*distributed as*,” while with three arguments they denote the corresponding pdf.

1.4.3 Probability Mass Function

The **probability mass function (pmf)** of a scalar random variable x , which can take values in the set $\{\xi_i, i = 1, \dots, n\}$ (i.e., it is discrete-valued), is

$$\mu_x(\xi_i) = P\{x = \xi_i\} = \mu_i \quad i = 1, \dots, n \quad (1.4.3-1)$$

where μ_i are the **point masses**.

Similarly to (1.4.2-5), the requirement for a **proper pmf** is

$$\sum_{i=1}^n \mu_i = 1 \quad (1.4.3-2)$$

Using the **Dirac (impulse) delta function** defined by

$$\delta(x) = 0 \quad \forall x \neq 0 \quad (1.4.3-3)$$

and

$$\int_{-\infty}^{\infty} \delta(x) dx = 1 \quad (1.4.3-4)$$

one can write a pdf corresponding to (1.4.3-1) as

$$p(x) = \sum_{i=1}^n \mu_i \delta(x - \xi_i) \quad (1.4.3-5)$$

Note that the above satisfies the **normalization property** (1.4.2-5) of a pdf in that it integrates to unity.

The distribution corresponding to the above density, called **cumulative probability mass function (cpmf)**, has jumps at ξ_i — it is a staircase function: Its derivative is zero everywhere except at the jumps where it is an impulse function.

The expression of the cpmf can be written, in terms of the **unit step function** $1(\cdot)$, as

$$P\{x \leq \xi\} = \sum_{i=1}^n \mu_i 1(\xi - \xi_i) \quad (1.4.3-6)$$

Note

For a random variable with no point mass at ξ , one has

$$P\{x \leq \xi\} = P\{x < \xi\} \quad (1.4.3-7)$$

Example

The **Poisson pmf** with rate λ

$$P\{x = n\} = e^{-\lambda T} \frac{(\lambda T)^n}{n!} \quad n = 0, 1, \dots \quad (1.4.3-8)$$

describes the number of random points in an interval T can also be written as

$$p(x) = \sum_{n=0}^{\infty} e^{-\lambda T} \frac{(\lambda T)^n}{n!} \delta(x - n) \quad (1.4.3-9)$$

1.4.4 Mixed Random Variable and Mixed Probability-PDF

A **mixed random variable** or a **hybrid random variable** x is one which can take values in a continuous set X as well as over a discrete set of points $\{\xi_i, i = 1, \dots, n\}$. Such a random variable has a pdf of the form

$$p(x) = p_c(x) + \sum_{i=1}^n \mu_i \delta(x - \xi_i) \quad (1.4.4-1)$$

where $p_c(x)$ is the continuous part of the pdf and μ_i are the point masses.

Then

$$\int_{-\infty}^{\infty} p(x) dx = \int_{x \in X} p_c(x) dx + \sum_{i=1}^n \mu_i = 1 \quad (1.4.4-2)$$

The **joint probability-pdf of an event and a random variable** is defined as

$$P_{A,x}[A, \xi] = \lim_{d\xi \rightarrow 0} \frac{P\{A, \xi - d\xi < x \leq \xi\}}{d\xi} \quad (1.4.4-3)$$

and is denoted without subscripts as $P[A, x]$ where the arguments define the function.

The following notations will be observed⁶ in the sequel:

$$P\{\cdot\} = \text{probability of an event} \quad (1.4.4-4)$$

$$p(\cdot) \text{ or } p[\cdot] = \text{pdf} \quad (1.4.4-5)$$

$$\mu(\cdot) = \text{pmf} \quad (1.4.4-6)$$

$$P[\cdot, \cdot] = \text{mixed (joint) probability-pdf} \quad (1.4.4-7)$$

Some of the uncertainties in estimation problems are naturally modeled as mixed random variables.

⁶Most of the time. Some of the time we will be guilty of using “dynamic notation” — the same variable keeps being denoted by new symbols, and the same symbol is used for different variables.

1.4.5 Expectations and Moments of a Scalar Random Variable

The **expected value** of a scalar random variable, also called its **mean**, **average**, or **first moment**, is

$$E[x] = \int_{-\infty}^{\infty} xp(x) dx \triangleq \bar{x} \quad (1.4.5-1)$$

The *n*th **moment** is

$$E[x^n] = \int_{-\infty}^{\infty} x^n p(x) dx \quad (1.4.5-2)$$

The **second central moment** or **variance** is

$$\text{var}(x) \triangleq E[(x - \bar{x})^2] = \int_{-\infty}^{\infty} (x - \bar{x})^2 p(x) dx = E[x^2] - (\bar{x})^2 \triangleq \sigma_x^2 \quad (1.4.5-3)$$

The square root σ of the variance is called the **standard deviation**.

A random variable x with mean \bar{x} and variance σ^2 will be denoted as

$$x \sim [\bar{x}, \sigma^2] \quad (1.4.5-4)$$

Note that this is similar to (1.4.2-10) except that the distribution is not specified.

The second moment, which is the **mean square (MS)** value, is equal to the square of the mean plus the variance

$$E[x^2] = [E(x)]^2 + \text{var}(x) = \bar{x}^2 + \sigma_x^2 \quad (1.4.5-5)$$

For a zero-mean random variable, the standard deviation is the **root mean square (RMS)** value.⁷

The expected value of a function $g(x)$ of the random variable x is

$$E[g(x)] = \int_{-\infty}^{\infty} g(x) p(x) dx \quad (1.4.5-6)$$

that is, the (Lebesgue) integral with respect to the measure $p(x) dx = dP(x)$.

Examples

If x is uniformly distributed in the interval $[a, b]$, that is,

$$p(x) = \mathcal{U}(x; a, b) \triangleq \begin{cases} \frac{1}{b-a} & x \in [a, b] \\ 0 & \text{elsewhere} \end{cases} \quad (1.4.5-7)$$

then one has

$$E[x] = \frac{b+a}{2} \quad \text{var}(x) = \frac{(b-a)^2}{12} \quad (1.4.5-8)$$

⁷Not to be confused with the RMS value of a waveform.

(See also problem 1-6.)

For a Gaussian random variable with parameters μ and σ^2 (see 1.4.2-10) the mean is

$$\begin{aligned} E[x] &= \int_{-\infty}^{\infty} x \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}} dx \\ &= \int_{-\infty}^{\infty} (x - \mu) \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}} dx + \mu \int_{-\infty}^{\infty} \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}} dx \\ &= \mu \end{aligned} \quad (1.4.5-9)$$

where in the last line above the first integral is zero due to symmetry and the second is unity due to the pdf normalization property.

The variance is

$$\begin{aligned} \text{var}(x) &= E[(x - \mu)^2] = \int_{-\infty}^{\infty} (x - \mu)^2 \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}} dx \\ &= \int_{-\infty}^{\infty} x^2 \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{x^2}{2\sigma^2}} dx \end{aligned} \quad (1.4.5-10)$$

The last integral above can be evaluated by differentiating w.r.t. σ the identity

$$\int_{-\infty}^{\infty} \frac{1}{\sqrt{2\pi}} e^{-\frac{x^2}{2\sigma^2}} dx = \sigma \quad (1.4.5-11)$$

which follows from the pdf normalization property. Thus

$$\int_{-\infty}^{\infty} \frac{1}{\sqrt{2\pi}\sigma^3} x^2 e^{-\frac{x^2}{2\sigma^2}} dx = 1 \quad (1.4.5-12)$$

Using this in (1.4.5-10) yields

$$\text{var}(x) = \sigma^2 \quad (1.4.5-13)$$

The mean of the Poisson random variable x from (1.4.3-9) is

$$E[x] = \int_{-\infty}^{\infty} x \sum_{n=0}^{\infty} e^{-\lambda T} \frac{(\lambda T)^n}{n!} \delta(x-n) dx = \sum_{n=0}^{\infty} n e^{-\lambda T} \frac{(\lambda T)^n}{n!} = \lambda T \quad (1.4.5-14)$$

1.4.6 Joint PDF of Two Random Variables

The **joint pdf** of two random variables x and y is defined in terms of the probability of the following joint event, denoted with the set intersection symbol (see Fig. 1.4.6-1), namely,

$$p_{x,y}(\xi, \eta) \triangleq \lim_{d\xi \rightarrow 0, d\eta \rightarrow 0} \frac{P\{\{\xi - d\xi < x \leq \xi\} \cap \{\eta - d\eta < y \leq \eta\}\}}{d\xi d\eta} \quad (1.4.6-1)$$

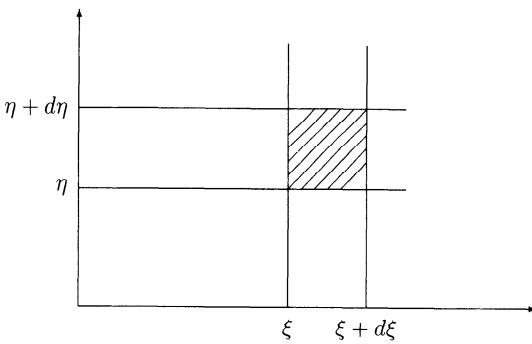


Figure 1.4.6-1: Event for definition of a joint pdf.

Integrating the joint pdf of two random variables over one of the variables yields the pdf of the other random variable

$$\int_{-\infty}^{\infty} p_{x,y}(\xi, \eta) d\eta = p_x(\xi) \quad (1.4.6-2)$$

or, using the simpler notation (1.4.2-2),

$$\int_{-\infty}^{\infty} p(x, y) dy = p(x) \quad (1.4.6-3)$$

The resulting pdf, which pertains to a single random variable (introduced in Subsection 1.4.2), is also called **marginal pdf** or **marginal density**.

Similarly to (1.4.2-4), the **joint cdf** is

$$P_{x,y}(\xi, \eta) = P\{x \leq \xi, y \leq \eta\} = \int_{x=-\infty}^{\xi} \int_{y=-\infty}^{\eta} p_{x,y}(x, y) dx dy \quad (1.4.6-4)$$

Covariance and Correlation Coefficient

The **covariance of two scalar random variables** x_1 and x_2 with means \bar{x}_1 and \bar{x}_2 , respectively, is

$$\begin{aligned} \text{cov}(x_1, x_2) &\triangleq E[(x_1 - \bar{x}_1)(x_2 - \bar{x}_2)] \\ &= \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} (x_1 - \bar{x}_1)(x_2 - \bar{x}_2) p(x_1, x_2) dx_1 dx_2 \triangleq \sigma_{x_1 x_2}^2 \end{aligned} \quad (1.4.6-5)$$

The **correlation coefficient** of these two random variables is the normalized quantity

$$\rho_{12} \triangleq \frac{\sigma_{x_1 x_2}^2}{\sigma_{x_1} \sigma_{x_2}} \quad (1.4.6-6)$$

where σ_{x_i} is the standard deviation of x_i .

Due to the normalization, the *magnitude of the correlation coefficient* of any two random variables obeys the following inequality:

$$|\rho_{12}| \leq 1 \quad (1.4.6-7)$$

It can be shown that this is a consequence of the Schwarz inequality (1.3.4-2) for random variables (where the inner product is the covariance and the norm squared is the variance — this is discussed in Section 3.3).

Two random variables whose correlation coefficient is zero are said to be **uncorrelated**. Clearly, the uncorrelatedness of two random variables is equivalent to their covariance being zero. Furthermore, the random variables x_1 and x_2 are uncorrelated if and only if

$$E[x_1 x_2] = E[x_1] E[x_2] \quad (1.4.6-8)$$

At the other extreme, if the correlation coefficient of two random variables has magnitude unity, then it can be shown that they are linearly dependent (i.e., one is a linear function of the other). Two random variables x_1 and x_2 are said to be **linearly dependent** if

$$a_1 x_1 + a_2 x_2 = 0 \quad (1.4.6-9)$$

for some $a = [a_1 \ a_2]' \neq 0$.

Example

Consider the random variables x and y with means \bar{x} , \bar{y} and variances σ_x^2 , σ_y^2 , respectively, and uncorrelated from each other (i.e., $\text{cov}(x, y) = 0$).

Let

$$z = ax + by \quad (1.4.6-10)$$

The mean of z is

$$\bar{z} = E[z] = E[ax + by] = aE[x] + bE[y] = a\bar{x} + b\bar{y} \quad (1.4.6-11)$$

while its variance is

$$\begin{aligned} \text{var}(z) &= E[(z - \bar{z})^2] = E[(ax + by - a\bar{x} - b\bar{y})^2] \\ &= E\{[a(x - \bar{x}) + b(y - \bar{y})]^2\} \\ &= a^2 E[(x - \bar{x})^2] + b^2 E[(y - \bar{y})^2] + 2ab E[(x - \bar{x})(y - \bar{y})] \\ &= a^2 \sigma_x^2 + b^2 \sigma_y^2 \end{aligned} \quad (1.4.6-12)$$

since the cross-term — the covariance of x and y — is zero.

The covariance of z and x is

$$\begin{aligned} E[(z - \bar{z})(x - \bar{x})] &= E\{[a(x - \bar{x}) + b(y - \bar{y})](x - \bar{x})\} \\ &= a\sigma_x^2 + b\text{cov}(x, y) = a\sigma_x^2 \end{aligned} \quad (1.4.6-13)$$

1.4.7 Independent Events and Independent Random Variables

Two events are *independent* if the probability of their joint event equals the product of their marginal probabilities

$$P\{A \cap B\} \triangleq P\{A, B\} = P\{A\}P\{B\} \quad (1.4.7-1)$$

A set of n events A_i , $i = 1, \dots, n$ are independent if the (joint) probability of their intersection is equal to the product of the corresponding (marginal) event probabilities

$$P\left\{\bigcap_{i=1}^n A_i\right\} = \prod_{i=1}^n P\{A_i\} \quad (1.4.7-2)$$

and the same property holds also for any subset of these events.

Similarly, n random variables are *independent* if their joint pdf equals the product of their corresponding marginal densities

$$p(x_1, \dots, x_n) = \prod_{i=1}^n p(x_i) \quad (1.4.7-3)$$

A set of random variables is called *independent, identically distributed (i.i.d.)* if (1.4.7-3) holds and their marginal distributions (or densities) are identical.

The pdf of the sum of two independent random variables

$$y = x_1 + x_2 \quad (1.4.7-4)$$

is the *convolution* of their marginal densities

$$p(y) = \int p_{x_1}(y - x)p_{x_2}(x) dx \quad (1.4.7-5)$$

(See also problem 1-1.)

1.4.8 Vector-Valued Random Variables and Their Moments

The pdf of the *vector-valued random variable*

$$x = [x_1 \ \cdots \ x_n]' \quad (1.4.8-1)$$

at

$$\xi = [\xi_1 \ \cdots \ \xi_n]' \quad (1.4.8-2)$$

is defined as the *joint density* of its components

$$p_{x_1, \dots, x_n}(\xi_1, \dots, \xi_n) \triangleq p_x(\xi) \triangleq \lim_{d\xi_1 \rightarrow 0, \dots, d\xi_n \rightarrow 0} \frac{P\{\bigcap_{i=1}^n \{\xi_i - d\xi_i < x_i \leq \xi_i\}\}}{d\xi_1 \cdots d\xi_n} \quad (1.4.8-3)$$

where the set intersection symbol \cap is used to denote a joint event

$$A \cap B \triangleq \{A \text{ and } B\} \triangleq \{A, B\} \quad (1.4.8-4)$$

The shorter notation $p(x)$ will be used if it does not lead to ambiguities.

The mean of the vector x is the result of the n -fold integration

$$E[x] = \int_{-\infty}^{\infty} \cdots \int_{-\infty}^{\infty} \xi p_x(\xi) d\xi_1 \cdots d\xi_n = \int_{-\infty}^{\infty} \cdots \int_{-\infty}^{\infty} xp(x) dx_1 \cdots dx_n \triangleq \bar{x} \quad (1.4.8-5)$$

The **covariance matrix** of the n -vector x is obtained from the n -fold integration (written with the short notation and with the limits omitted)

$$\text{cov}(x) \triangleq E[(x - \bar{x})(x - \bar{x})'] = \int (x - \bar{x})(x - \bar{x})' p(x) dx \triangleq P_{xx} \quad (1.4.8-6)$$

where $dx \triangleq \prod dx_i$. Note that the covariance matrix is a *symmetric* $n \times n$ matrix.

The diagonal elements of the *covariance matrix* are the *variances* of the components of x , while the off-diagonal elements are the (*scalar*) *covariances* between its components, as in (1.4.6-5).

The *covariance matrix* is *positive definite* (and thus nonsingular) unless there is a linear dependence among the components of x . If there is such a dependence, the covariance matrix is positive semidefinite.

Note

Covariance matrices will be denoted by the symbol P , with a subscript or argument (in parentheses); the designation of a probability will be $P\{\cdot\}$ where braces are used for the event it pertains to, as in (1.4.4-4).

Characteristic Function and Moments

The **characteristic function** of a vector random variable is defined as the n -fold integral (with the limits, which are, as in (1.4.8-5), omitted)

$$M_x(s) = E[e^{s'x}] = \int e^{s'x} p(x) dx \quad (1.4.8-7)$$

which is the (n -dimensional) Fourier transform of the pdf with argument the (purely imaginary) vector s of the same dimension n as that of x .

The first moment of x can be obtained from the characteristic function as

$$E[x] = \nabla_s M_x(s)|_{s=0} \quad (1.4.8-8)$$

that is, its gradient is evaluated at $s = 0$, where ∇_s is the (column) gradient operator

$$\nabla_s = \left[\frac{\partial}{\partial s_1} \cdots \frac{\partial}{\partial s_n} \right]' \quad (1.4.8-9)$$

Similarly,

$$E[xx'] = \nabla_s \nabla'_s M_x(s)|_{s=0} \quad (1.4.8-10)$$

and so forth.

Due to this property, the characteristic function is also called the **moment-generating function**.

Example

The characteristic function of the scalar Gaussian random variable

$$x \sim \mathcal{N}(\mu, \sigma^2) \quad (1.4.8-11)$$

is

$$M_x(s) = e^{s\mu + \frac{1}{2}s^2\sigma^2} \quad (1.4.8-12)$$

To find the moments of x it is convenient to deal with the zero-mean random variable

$$y = x - \mu \quad (1.4.8-13)$$

The moments of x follow easily from the moments of y , according to

$$E[x^n] = E[(y + \mu)^n] \quad (1.4.8-14)$$

The characteristic function of y is

$$M_y(s) = e^{\frac{1}{2}s^2\sigma^2} \quad (1.4.8-15)$$

and

$$\frac{d}{ds} M_y(s) = s\sigma^2 e^{\frac{1}{2}s^2\sigma^2} \quad (1.4.8-16)$$

Evaluating the above at $s = 0$ yields the first moment (mean) of y as 0.

Similarly,

$$\frac{d^2}{ds^2} M_y(s) = [\sigma^2 + (s\sigma^2)^2] e^{\frac{1}{2}s^2\sigma^2} \quad (1.4.8-17)$$

which, at $s = 0$, yields

$$E[y^2] = E[(x - \mu)^2] = \text{var}(x) = \sigma^2 \quad (1.4.8-18)$$

The third derivative of the characteristic function of y is

$$\frac{d^3}{ds^3} M_y(s) = [2s\sigma^4 + (\sigma^2 + s^2\sigma^4)s\sigma^2] e^{\frac{1}{2}s^2\sigma^2} = [3s\sigma^4 + s^3\sigma^6] e^{\frac{1}{2}s^2\sigma^2} \quad (1.4.8-19)$$

which yields zero as the third moment of y .

The fourth moment of y is given by

$$\frac{d^4}{ds^4} M_y(s) = [3\sigma^4 + 3s^2\sigma^6 + (3s\sigma^4 + s^3\sigma^6)s\sigma^2] e^{\frac{1}{2}s^2\sigma^2} \quad (1.4.8-20)$$

at $s = 0$ as

$$E[y^4] = E[(x - \mu)^4] = 3\sigma^4 \quad (1.4.8-21)$$

1.4.9 Conditional Probability and PDF

The **conditional probability** of an event A given B is defined as

$$P\{A|B\} = \frac{P\{A, B\}}{P\{B\}} \quad (1.4.9-1)$$

For independent events, the above becomes the unconditional probability.

Similarly, the **conditional pdf** of one random variable given another random variable is, using the simpler notation as in (1.4.2-2), given by

$$p(x|y) = \frac{p(x, y)}{p(y)} \quad (1.4.9-2)$$

For an event conditioned on a random variable, one has

$$P\{A|x\} = \frac{P[A, x]}{p(x)} \quad (1.4.9-3)$$

The conditional pdf of a random variable x , given an event A , is

$$p(x|A) = \frac{P[A, x]}{P\{A\}} \quad (1.4.9-4)$$

where the numerator above is a mixed probability-pdf, as defined in (1.4.4-3).

Example

If the conditioning event is

$$A = \{x \leq a\} \quad (1.4.9-5)$$

then one gets the **truncated pdf**

$$p(x|x \leq a) = \frac{P[x, x \leq a]}{P\{x \leq a\}} = \begin{cases} \frac{p(x)}{P\{x \leq a\}} & \text{if } x \leq a \\ 0 & \text{elsewhere} \end{cases} \quad (1.4.9-6)$$

which is the original one restricted to $x \leq a$ and suitably *renormalized*.

If $z = x + y$ and the conditional pdf of y given x is $p_{y|x}(\cdot)$, then

$$p(z|x) = p_{y|x}(z - x) \quad (1.4.9-7)$$

If, in addition, x and y are independent, then, with $p_y(\cdot)$ denoting the pdf of y , one has

$$p(z|x) = p_y(z - x) \quad (1.4.9-8)$$

(See also problem 1-1.)

Assume $x \sim \mathcal{N}(\mu_x, \sigma_x^2)$, $y \sim \mathcal{N}(\mu_y, \sigma_y^2)$ and independent. Then, with $z = x + y$,

$$\begin{aligned} p(z|x) &= p_y(z-x) = \mathcal{N}(z-x; \mu_y, \sigma_y^2) \\ &= \frac{1}{\sqrt{2\pi}\sigma_y} e^{-\frac{(z-x-\mu_y)^2}{2\sigma_y^2}} = \mathcal{N}(z; x+\mu_y, \sigma_y^2) \end{aligned} \quad (1.4.9-9)$$

Note that the above can be viewed, due to the special form of the Gaussian density, as the pdf of $z - x$ with mean μ_y or the pdf of z with mean $x + \mu_y$.

If $x \sim \mathcal{N}(\mu, \sigma^2)$, then

$$p(x| |x - \mu| \leq a) = \begin{cases} \frac{1}{P\{|x - \mu| \leq a\}} \mathcal{N}(x; \mu, \sigma^2) & |x| \leq a \\ 0 & \text{elsewhere} \end{cases} \quad (1.4.9-10)$$

is a **truncated Gaussian**.

The normalization constant above is the probability that $x - \mu$ lies in the closed interval⁸ $[-a, a]$, which is

$$\begin{aligned} P\{|x - \mu| \leq a\} &= P\{x - \mu \in [-a, a]\} = P\{x \in [\mu - a, \mu + a]\} \\ &= P\{x \in (-\infty, \mu + a]\} - P\{x \in (-\infty, \mu - a)\} \\ &\equiv \mathcal{G}\left(\frac{(\mu + a) - \mu}{\sigma}\right) - \mathcal{G}\left(\frac{(\mu - a) - \mu}{\sigma}\right) \\ &= \mathcal{G}\left(\frac{a}{\sigma}\right) - \mathcal{G}\left(\frac{-a}{\sigma}\right) = 1 - 2\mathcal{G}\left(\frac{-a}{\sigma}\right) \end{aligned} \quad (1.4.9-11)$$

where \mathcal{G} is the cumulative standard Gaussian distribution and the last equality follows from the symmetry of the Gaussian density. The last term above is the tail probability to the left of $(-a)/\sigma$ or to the right of $(a)/\sigma$.

1.4.10 The Total Probability Theorem

Let the events B_i , $i = 1, \dots, n$, be **mutually exclusive**, that is,

$$P\{B_i, B_j\} = 0 \quad \forall i \neq j \quad (1.4.10-1)$$

and **exhaustive**, that is,

$$\sum_{i=1}^n P\{B_i\} = 1 \quad (1.4.10-2)$$

Such a set of events is a *partition of the space of outcomes*. Then, the **total probability theorem** states that for any event A , its probability can be decomposed in terms of conditional probabilities as follows:

$$P\{A\} = \sum_{i=1}^n P\{A, B_i\} = \sum_{i=1}^n P\{A|B_i\}P\{B_i\} \quad (1.4.10-3)$$

⁸A bracket indicates that the interval is closed at the corresponding end, i.e., it includes the endpoint. A parenthesis indicates that the interval is open at the corresponding end, i.e., it excludes the endpoint.

The key to the above is the mutual exclusiveness and exhaustiveness of the set of events B_i , $i = 1, \dots, n$.

Figure 1.4.10-1 illustrates this with the set theory counterpart, namely,

$$A = \bigcup_{i=1}^n (A \cap B_i) \quad (1.4.10-4)$$

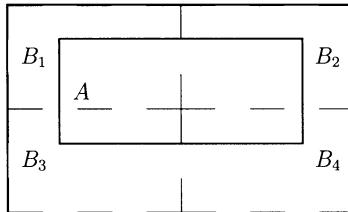


Figure 1.4.10-1: The total probability theorem (set theory counterpart).

The version of the total probability theorem for random variables is

$$p(x) = \int_{-\infty}^{\infty} p(x, y) dy = \int_{-\infty}^{\infty} p(x|y)p(y) dy \quad (1.4.10-5)$$

Note that this is a combination of the definition of the conditional pdf (1.4.9-2) and property (1.4.6-3) that the joint pdf integrated out with respect to one random variable yields the marginal pdf of the other variable.

For mixed event/random variable situations, one has

$$p(x) = \sum_{i=1}^n p(x|B_i)P\{B_i\} \quad (1.4.10-6)$$

(this is illustrated in an example in Subsection 1.4.14) and

$$P\{A\} = \int_{-\infty}^{\infty} P\{A|x\}p(x) dx = E[P\{A|x\}] \quad (1.4.10-7)$$

where the expectation operator averages over x .

An additional *common* conditioning in *all* the probabilities is permissible, for instance,

$$P\{A|C\} = \sum_{i=1}^n P\{A, B_i|C\} = \sum_{i=1}^n P\{A|B_i, C\}P\{B_i|C\} \quad (1.4.10-8)$$

or

$$P\{A|y\} = \int_{-\infty}^{\infty} P\{A|x, y\}p(x|y) dx \quad (1.4.10-9)$$

Similarly,

$$p(x|z) = \int_{-\infty}^{\infty} p(x, y|z) dy = \int_{-\infty}^{\infty} p(x|y, z)p(y|z) dy \quad (1.4.10-10)$$

The total probability theorem is the primary tool in obtaining the state estimate in the presence of extraneous uncertainties (e.g., model uncertainties or measurements of uncertain origin).

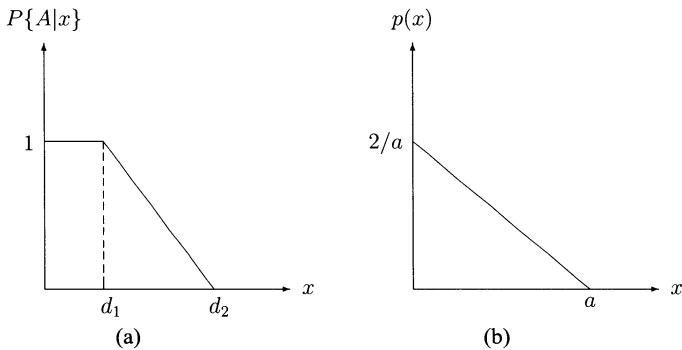


Figure 1.4.10-2: Success probability and miss distance pdf.

Example

The probability of event A , which denotes success in a hunting target shooting, can be modeled as a function of the “miss distance” x by the conditional probability (see Fig. 1.4.10-2a)

$$P\{A|x\} = \begin{cases} 1 & x \leq d_1 \\ \frac{x - d_2}{d_1 - d_2} & d_1 < x \leq d_2 \\ 0 & x > d_2 \end{cases} \quad (1.4.10-11)$$

The miss distance pdf can be taken as (see Fig. 1.4.10-2b)

$$p(x) = \begin{cases} \frac{2(a-x)}{a^2} & 0 < x \leq a \\ 0 & \text{otherwise} \end{cases} \quad (1.4.10-12)$$

Then, assuming $d_2 < a$, the (unconditional) probability of success is

$$\begin{aligned} P\{A\} &= \int_{-\infty}^{\infty} P\{A|x\}p(x) dx = \int_0^{d_1} p(x) dx + \int_{d_1}^{d_2} \frac{x - d_2}{d_1 - d_2} p(x) dx \\ &= P\{x \leq d_1\} + \int_{d_1}^{d_2} \frac{x - d_2}{d_1 - d_2} \frac{2(a-x)}{a^2} dx \\ &= \frac{d_1(2a - d_1)}{a^2} + \frac{(d_2 - d_1)}{3a^2} (3a - 2d_1 - d_2) \end{aligned} \quad (1.4.10-13)$$

1.4.11 Bayes' Formula

The probability of an event B_i conditioned on event A can be expressed in terms of the reverse conditioning as follows:

$$P\{B_i|A\} = \frac{P\{A|B_i\}P\{B_i\}}{P\{A\}} \quad (1.4.11-1)$$

This is known as **Bayes' formula** or **Bayes' theorem** and is also referred to as **Bayes' rule**.⁹

The conditional probability of B_i is sometimes referred to as **posterior probability**, while the unconditional one is referred to as **prior probability**.

Using the total probability theorem (1.4.10-3) for the denominator in the above, one also has (if B_j , $j = 1, \dots, n$ are mutually exclusive and exhaustive)

$$P\{B_i|A\} = \frac{P\{A|B_i\}P\{B_i\}}{\sum_{j=1}^n P\{A|B_j\}P\{B_j\}} = \frac{1}{c}P\{A|B_i\}P\{B_i\} \quad (1.4.11-2)$$

where now the denominator c appears clearly as the **normalizing constant**, which guarantees that

$$\sum_{i=1}^n P\{B_i|A\} = 1 \quad (1.4.11-3)$$

Bayes' formula for random variables is written as

$$p(x|y) = \frac{p(y|x)p(x)}{p(y)} = \frac{p(y|x)p(x)}{\int p(y|x)p(x) dx} \quad (1.4.11-4)$$

In this case the unconditional pdf $p(x)$ is also called the **prior pdf** and the conditional pdf $p(x|y)$ is also called the **posterior pdf**.

For a mixed case

$$P\{B_i|x\} = \frac{p(x|B_i)P\{B_i\}}{p(x)} = \frac{p(x|B_i)P\{B_i\}}{\sum_{j=1}^n p(x|B_j)P\{B_j\}} \quad (1.4.11-5)$$

where the denominator is again the normalizing factor.

When there are several conditioning random variables or events, Bayes' formula can be used to "switch" only some of them. For example,

$$P\{B_i|x, y\} = \frac{p(x|B_i, y)P\{B_i|y\}}{p(x|y)} = \frac{p(x|B_i, y)P\{B_i|y\}}{\sum_{j=1}^n p(x|B_j, y)P\{B_j|y\}} \quad (1.4.11-6)$$

These equations are the key tools in state estimation for **hybrid systems** — systems with continuous as well as discrete uncertainties.

Remarks

The prior $p(x)$ reflects the (possibly subjective¹⁰) initial **degree of belief**, which, when combined with the **evidence from the data** — the **likelihood function** of x , namely, $p(y|x)$ — yields the posterior $p(x|y)$.

⁹**Bayes' decision rule** (or **Bayes' criterion** or **Bayes' principle**) in a decision problem is to minimize the expected value of a cost function, called **Bayes' risk**.

¹⁰Similarly to the well-known fact that the entire population of the world is divided into two groups (those who believe this and those who don't), the people dealing with probability are divided into two groups: proud Bayesians (who believe in priors) and closet Bayesians.

Bayes' postulate states that in the absence of prior knowledge, a **uniform prior pdf** should be chosen. While this is reasonable for discrete valued random variables, for continuous valued random variables this would imply a nonuniform pdf on any nonlinear function of this random variable and this has generated a lot of controversy among statisticians.

The use of a uniform prior pdf *over an infinite interval*, called **diffuse** or **noninformative**, also called **Dirichlet prior** is discussed in more detail in Subsection 2.3.4.

Example

The “noise” in a device, denoted as x and modeled as a random variable, is observed. If the device is operating according to its specifications, then one has event B_1 , defined as

$$B_1 = \{x \sim \mathcal{N}(0, \sigma_1^2)\} \quad (1.4.11-7)$$

The defective device event is assumed for simplicity defined as

$$B_2 = \{x \sim \mathcal{N}(0, \sigma_2^2)\} \quad (1.4.11-8)$$

The conditional probability that the device is defective given the observation x is

$$P\{B_2|x\} = \frac{p(x|B_2)P\{B_2\}}{p(x|B_1)P\{B_1\} + p(x|B_2)P\{B_2\}} \quad (1.4.11-9)$$

Using the explicit expressions of the Gaussian densities defining the events B_1 and B_2 , the above can be written, after cancellations, as

$$P\{B_2|x\} = \frac{1}{1 + \frac{\sigma_2}{\sigma_1} e^{-\frac{x^2}{2\sigma_1^2} + \frac{x^2}{2\sigma_2^2} \frac{P\{B_1\}}{P\{B_2\}}}} \quad (1.4.11-10)$$

Assume that a priori the two events are equiprobable

$$P\{B_1\} = \frac{1}{2} \quad P\{B_2\} = 1 - P\{B_1\} = \frac{1}{2} \quad (1.4.11-11)$$

If $\sigma_2 = 4\sigma_1$ and the observation obtained is $x = \sigma_2$, then

$$P\{B_2|x\} = \frac{1}{1 + 4e^{-8+0.5}} \approx 0.998 \quad (1.4.11-12)$$

If the prior is different, namely, $P\{B_2\} = 10^{-2}$, then

$$P\{B_2|x\} = \frac{1}{1 + 396e^{-8+0.5}} \approx 0.82 \quad (1.4.11-13)$$

i.e., regardless of the prior, the “evidence from the data” indicates a very high probability that the device is defective.

1.4.12 Conditional Expectations and Their Smoothing Property

The **conditional expectation** is defined similarly to the unconditional expectation (1.4.5-1) but *with respect to a conditional pdf*, that is,

$$E[x|y] = \int_{-\infty}^{\infty} xp(x|y) dx \quad (1.4.12-1)$$

Similarly, for a function of the random variable x (and possibly y), one has

$$E[g(x, y)|y] = \int_{-\infty}^{\infty} g(x, y)p(x|y) dx \quad (1.4.12-2)$$

Note that (1.4.12-1) and (1.4.12-2) are functions of the conditioning argument y .

The **smoothing property of the expectations** states that the expected value of a conditional expected value is the (unconditional) expected value:

$$\begin{aligned} E\{E[x|y]\} &= \int_{y=-\infty}^{\infty} \left[\int_{x=-\infty}^{\infty} xp(x|y) dx \right] p(y) dy \\ &= \int_{x=-\infty}^{\infty} x \left[\int_{y=-\infty}^{\infty} p(x, y) dy \right] dx \\ &= \int_{x=-\infty}^{\infty} xp(x) dx = E[x] \end{aligned} \quad (1.4.12-3)$$

In the first line of the above, the inside expectation is a function of y , which is “averaged out” (integrated over) by the outside expectation.

Equation (1.4.12-3) is also called the **law of iterated expectations**, summarized as

$$\boxed{E[E[x|y]] = E[x]} \quad (1.4.12-4)$$

The same property holds when the conditioning is on events or mixed random variables.

This is used in the evaluation of the performance of certain estimation algorithms operating in the presence of several types of uncertainties — it allows us to handle different types of uncertainties sequentially rather than simultaneously.

Example

To illustrate the smoothing property of expectations, consider the conditional mean

$$E(x|z) = \bar{x} + P_{xz}P_{zz}^{-1}(z - \bar{z}) \quad (1.4.12-5)$$

where x and z are jointly Gaussian (this is derived later in (1.4.14-17)); overbar denotes the unconditional expected values.

Applying the unconditional expectation operator on the above, where the only random variable on the right-hand side is z , with mean \bar{z} , yields

$$E[E(x|z)] = E[\bar{x}] + P_{xz}P_{zz}^{-1} E(z - \bar{z}) = \bar{x} = E[x] \quad (1.4.12-6)$$

1.4.13 Gaussian Random Variables

The pdf of a (scalar) **Gaussian** or **normal random variable** is

$$p(x) = \mathcal{N}(x; \bar{x}, \sigma^2) \triangleq \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\bar{x})^2}{2\sigma^2}} \quad (1.4.13-1)$$

where $\mathcal{N}(\cdot)$ denotes the **normal pdf (density)** with argument x , mean \bar{x} , and variance σ^2 . The first two moments, which fully characterize a Gaussian random variable, are referred to as its **statistics**.

Another notation equivalent to the above is

$$x \sim \mathcal{N}(\bar{x}, \sigma^2) \quad (1.4.13-2)$$

which states that x is *normally distributed* with the corresponding mean and variance. Note the different meanings of \mathcal{N} in (1.4.13-1) and (1.4.13-2) — they are, however, specified by their arguments: three in (1.4.13-1) vs. two in (1.4.13-2).

A **vector-valued Gaussian** random variable has the density

$$\mathcal{N}(x; \bar{x}, P) \triangleq |2\pi P|^{-1/2} e^{-\frac{1}{2}(x-\bar{x})' P^{-1}(x-\bar{x})} \quad (1.4.13-3)$$

where

$$\bar{x} = E[x] \quad (1.4.13-4)$$

$$P = E[(x - \bar{x})(x - \bar{x})'] \quad (1.4.13-5)$$

are, respectively, the mean and covariance matrix of the vector x .

The determinant in (1.4.13-3) has been written with the factor 2π inside it by making use of (1.3.3-5). This avoids the need to indicate the dimension of the vector x in (1.4.13-3).

The components of a Gaussian distributed vector are said to be **jointly Gaussian**.

If the covariance matrix P is diagonal — that is, the components of the **Gaussian** random vector x are *uncorrelated* — then they are also *independent* because their joint pdf equals the product of the marginals.

A very important property of Gaussian random variables is that they *stay Gaussian under linear transformations*. Consequently, the sum of (jointly) Gaussian random variables is also Gaussian. This property is a major reason that the Gaussian model is very commonly used in estimation.

Example

If x_1 and x_2 are random vectors of the same dimension and jointly Gaussian distributed, that is,

$$p\left(\begin{bmatrix} x_1 \\ x_2 \end{bmatrix}\right) = \mathcal{N}\left\{\begin{bmatrix} x_1 \\ x_2 \end{bmatrix}; \begin{bmatrix} \bar{x}_1 \\ \bar{x}_2 \end{bmatrix}, \begin{bmatrix} P_{11} & P_{12} \\ P_{21} & P_{22} \end{bmatrix}\right\} \quad (1.4.13-6)$$

then

$$x = x_1 + x_2 \quad (1.4.13-7)$$

is Gaussian

$$p(x) = \mathcal{N}(x; \bar{x}, P) \quad (1.4.13-8)$$

where

$$\bar{x} = \bar{x}_1 + \bar{x}_2 \quad (1.4.13-9)$$

$$P = P_{11} + P_{12} + P_{21} + P_{22} \quad (1.4.13-10)$$

Symbolically, one can write, with notation (1.4.13-2),

$$\mathcal{N}(\bar{x}_1, P_{11}) + \mathcal{N}(\bar{x}_2, P_{22}) = \mathcal{N}(\bar{x}_1 + \bar{x}_2, P_{11} + P_{12} + P_{21} + P_{22}) \quad (1.4.13-11)$$

If x_1 and x_2 independent, the pdf of their sum $x = x_1 + x_2$ conditioned on x_1 is

$$\begin{aligned} p(x|x_1) &= p_{x_2}(x - x_1) \\ &= \mathcal{N}(x - x_1; \bar{x}_2, P_{22}) = |2\pi P_{22}|^{-1/2} e^{-\frac{1}{2}(x-x_1-\bar{x}_2)'P_{22}^{-1}(x-x_1-\bar{x}_2)} \\ &= \mathcal{N}(x; x_1 + \bar{x}_2, P_{22}) \end{aligned} \quad (1.4.13-12)$$

Note that in (1.4.13-12) x_1 acts as a constant — it shifts the mean — since it is given in the conditioning.

1.4.14 Joint and Conditional Gaussian Random Variables

Two random vectors x and z are *jointly Gaussian* if the *stacked vector*

$$y = \begin{bmatrix} x \\ z \end{bmatrix} \quad (1.4.14-1)$$

is Gaussian, that is,

$$p(x, z) = p(y) = \mathcal{N}(y; \bar{y}, P_{yy}) \quad (1.4.14-2)$$

The mean and covariance matrix of y in terms of those of x and z are

$$\bar{y} = \begin{bmatrix} \bar{x} \\ \bar{z} \end{bmatrix} \quad (1.4.14-3)$$

$$P_{yy} = \begin{bmatrix} P_{xx} & P_{xz} \\ P_{zx} & P_{zz} \end{bmatrix} \quad (1.4.14-4)$$

where

$$P_{xx} = \text{cov}(x) = E[(x - \bar{x})(x - \bar{x})'] \quad (1.4.14-5)$$

$$P_{zz} = \text{cov}(z) = E[(z - \bar{z})(z - \bar{z})'] \quad (1.4.14-6)$$

$$P_{xz} = \text{cov}(x, z) = E[(x - \bar{x})(z - \bar{z})'] = P'_{zx} \quad (1.4.14-7)$$

are the blocks of the **partitioned covariance matrix** (1.4.14-4).

The conditional pdf of x given z when they are jointly Gaussian is

$$p(x|z) = \frac{p(x, z)}{p(z)} = \frac{|2\pi P_{yy}|^{-1/2} e^{-\frac{1}{2}(y - \bar{y})' P_{yy}^{-1}(y - \bar{y})}}{|2\pi P_{zz}|^{-1/2} e^{-\frac{1}{2}(z - \bar{z})' P_{zz}^{-1}(z - \bar{z})}} \quad (1.4.14-8)$$

Advantage will be taken of the fact that the above is an exponential whose exponent is the difference of the exponents of the numerator and denominator.

The following notations will be used

$$\xi \triangleq x - \bar{x} \quad (1.4.14-9)$$

$$\zeta \triangleq z - \bar{z} \quad (1.4.14-10)$$

Using the new variables ξ and ζ , the problem for the *nonzero mean random variables* x and y is reduced to that of *zero-mean variables*.

The exponent on the right-hand side of the conditional density (1.4.14-8) is (after multiplication by -2) the quadratic form

$$\begin{aligned} q &= \begin{bmatrix} \xi \\ \zeta \end{bmatrix}' \begin{bmatrix} P_{xx} & P_{xz} \\ P_{zx} & P_{zz} \end{bmatrix}^{-1} \begin{bmatrix} \xi \\ \zeta \end{bmatrix} - \zeta' P_{zz}^{-1} \zeta \\ &= \begin{bmatrix} \xi \\ \zeta \end{bmatrix}' \begin{bmatrix} T_{xx} & T_{xz} \\ T_{zx} & T_{zz} \end{bmatrix} \begin{bmatrix} \xi \\ \zeta \end{bmatrix} - \zeta' P_{zz}^{-1} \zeta \end{aligned} \quad (1.4.14-11)$$

The relationships between the partitions of the inverse of the covariance matrix and the partitions of the original matrix are, using (1.3.3-7) to (1.3.3-10), given by

$$T_{xx}^{-1} = P_{xx} - P_{xz} P_{zz}^{-1} P_{zx} \quad (1.4.14-12)$$

$$P_{zz}^{-1} = T_{zz} - T_{zx} T_{xx}^{-1} T_{xz} \quad (1.4.14-13)$$

$$T_{xx}^{-1} T_{xz} = -P_{xz} P_{zz}^{-1} \quad (1.4.14-14)$$

The exponent (1.4.14-11) can be rewritten as

$$\begin{aligned} q &= \xi' T_{xx} \xi + \xi' T_{xz} \zeta + \zeta' T_{zx} \xi + \zeta' T_{zz} \zeta - \zeta' P_{zz}^{-1} \zeta \\ &= (\xi + T_{xx}^{-1} T_{xz} \zeta)' T_{xx} (\xi + T_{xx}^{-1} T_{xz} \zeta) + \zeta' (T_{zz} - T_{zx} T_{xx}^{-1} T_{xz}) \zeta - \zeta' P_{zz}^{-1} \zeta \\ &= (\xi + T_{xx}^{-1} T_{xz} \zeta)' T_{xx} (\xi + T_{xx}^{-1} T_{xz} \zeta) \end{aligned} \quad (1.4.14-15)$$

where use has been made of (1.4.14-13). The above procedure is called **completion of the squares** (actually, of the quadratic forms).

The result is a quadratic form in x , meaning that the conditional pdf of x given z is also Gaussian. This can be seen as follows: In view of (1.4.14-9), (1.4.14-10), and (1.4.14-14), the expression on the right-hand side of (1.4.14-15) is a quadratic form in

$$\xi + T_{xx}^{-1} T_{xz} \zeta = x - \bar{x} - P_{xz} P_{zz}^{-1} (z - \bar{z}) \quad (1.4.14-16)$$

From this one can recognize the conditional mean of x given z as

$$E(x|z) \triangleq \hat{x} = \bar{x} + P_{xz}P_{zz}^{-1}(z - \bar{z}) \quad (1.4.14-17)$$

The corresponding conditional covariance is

$$\text{cov}(x|z) \triangleq P_{xx|z} = T_{xx}^{-1} = P_{xx} - P_{xz}P_{zz}^{-1}P_{zx} \quad (1.4.14-18)$$

Note that the conditional mean (1.4.14-17) is linear in the observation z and that the covariance (1.4.14-18) is independent of the observation.

The above are the **fundamental equations of linear estimation**.

1.4.15 Expected Value of Quadratic and Quartic Forms

Consider a vector-valued random variable x with mean $\bar{x} = 0$ and covariance matrix P . Then, the **expected value of a quadratic form** with this random vector can be written as

$$E[x'Ax] = E[\text{tr}(x'Ax)] = E[\text{tr}(Axx')] = \text{tr}[AE(xx')] = \text{tr}(AP) \quad (1.4.15-1)$$

The same result can be obtained for a Gaussian random vector using its characteristic function as follows.

If

$$x \sim \mathcal{N}(\bar{x}, P) \quad (1.4.15-2)$$

then its characteristic (or moment generating) function is

$$M_x(s) = E[e^{s'x}] = e^{\frac{1}{2}s'Ps+s'\bar{x}} \quad (1.4.15-3)$$

Using the properties of the gradient, one can write

$$\begin{aligned} E[x'Ax] &= E[(\nabla_s e^{s'x})'Ax]|_{s=0} = E[\nabla'_s Ax e^{s'x}]|_{s=0} \\ &= \nabla'_s A E[x e^{s'x}]|_{s=0} = \nabla'_s A E[\nabla_s e^{s'x}]|_{s=0} \\ &= \nabla'_s A \nabla_s M_x(s)|_{s=0} = \nabla'_s A \nabla_s M_x(s)|_{s=0} \end{aligned} \quad (1.4.15-4)$$

Using the characteristic function (1.4.15-3) with $\bar{x} = 0$ in the above yields (1.4.15-1) after the evaluation of the gradient. This technique, even though not the simplest for a quadratic form, can be used conveniently to evaluate the expected value of a quartic form, which is needed for the covariance of quadratic forms.

Note that the expected value of a quadratic form (1.4.15-1) holds regardless of the distribution of x .

For the **quartic form** to be considered next, the result will rely on the fact that x is Gaussian.

Analogously to (1.4.15-4), one has the **expected value of a quartic form** written as follows:

$$E[x'Ax x' Bx] = \nabla'_s A \nabla_s \nabla'_s B \nabla_s M_x(s)|_{s=0} \quad (1.4.15-5)$$

Using (1.4.15-3), it can be shown that after some computations the final result is

$$E[x'Ax x' Bx] = \text{tr}(AP)\text{tr}(BP) + 2\text{tr}(APBP) \quad (1.4.15-6)$$

The scalar version of the above for $A = B = 1$ and $P = \sigma^2$ is the well-known expression

$$E[x^4] = 3\sigma^4 \quad (1.4.15-7)$$

The **covariance of two quadratic forms** is, using (1.4.15-6),

$$\begin{aligned} E[[x'Ax - E(x'Ax)][x'Bx - E(x'Bx)]] &= E[x'Ax x' Bx] - E[x'Ax]E[x'Bx] \\ &= 2\text{tr}(APBP) \end{aligned} \quad (1.4.15-8)$$

These results are used in the state estimation for nonlinear systems where nonlinear functions (system dynamics or measurement equation) are approximated by a series expansion of up to second order.

1.4.16 Mixture Probability Density Functions

A **mixture pdf** is a **weighted sum of pdfs** with the weights summing up to unity.

A **Gaussian mixture** is a pdf consisting of a weighted sum of Gaussian densities

$$p(x) = \sum_{j=1}^n p_j \mathcal{N}(x; \bar{x}_j, P_j) \quad (1.4.16-1)$$

where, obviously,

$$\sum_{j=1}^n p_j = 1 \quad (1.4.16-2)$$

Denote by A_j the event that x is Gaussian distributed with mean \bar{x}_j and covariance P_j , that is,

$$A_j \triangleq \{x \sim \mathcal{N}(\bar{x}_j, P_j)\} \quad (1.4.16-3)$$

With A_j , $j = 1, \dots, n$, mutually exclusive and exhaustive, and

$$P\{A_j\} = p_j \quad (1.4.16-4)$$

one can rewrite (1.4.16-1) with the Total Probability Theorem as

$$p(x) = \sum_{j=1}^n p(x|A_j)P\{A_j\} \quad (1.4.16-5)$$

The Moments of a Mixture

The mean of a mixture is easily seen to be

$$\bar{x} = \sum_{j=1}^n p_j \bar{x}_j \quad (1.4.16-6)$$

that is, the weighted sum of the means of the component densities regardless of the latter's nature.

The covariance of this mixture is

$$\begin{aligned} E[(x - \bar{x})(x - \bar{x})'] &= \sum_{j=1}^n E[(x - \bar{x})(x - \bar{x})' | A_j] p_j \\ &= \sum E[(x - \bar{x}_j + \bar{x}_j - \bar{x})(x - \bar{x}_j + \bar{x}_j - \bar{x})' | A_j] p_j \\ &= \sum p_j E[(x - \bar{x}_j)(x - \bar{x}_j)' | A_j] + \sum (\bar{x}_j - \bar{x})(\bar{x}_j - \bar{x})' p_j \end{aligned} \quad (1.4.16-7)$$

which can be written as

$$E[(x - \bar{x})(x - \bar{x})'] = \sum_{j=1}^n p_j P_j + \tilde{P} \quad (1.4.16-8)$$

In the above

$$\tilde{P} \triangleq \sum (\bar{x}_j - \bar{x})(\bar{x}_j - \bar{x})' p_j \quad (1.4.16-9)$$

is the **spread of the means term**.

An alternative expression for the above is

$$\begin{aligned} \tilde{P} &= \sum \bar{x}_j \bar{x}'_j p_j - \bar{x} \sum \bar{x}'_j p_j - \sum \bar{x}_j p_j \bar{x}' + \bar{x} \bar{x}' \sum p_j \\ &= \sum p_j \bar{x}_j \bar{x}'_j - \bar{x} \bar{x}' \end{aligned} \quad (1.4.16-10)$$

which yields another form for (1.4.16-8)

$$E[(x - \bar{x})(x - \bar{x})'] = \sum_{j=1}^n p_j P_j + \sum_{j=1}^n p_j \bar{x}_j \bar{x}'_j - \bar{x} \bar{x}' \quad (1.4.16-11)$$

The spread of the means term \tilde{P} , defined in (1.4.16-9), is a sum of dyads with positive weightings, and thus it follows from Subsection 1.3.6 that \tilde{P} is **positive semidefinite**.

Note that (1.4.16-6) and (1.4.16-8), which will be referred to as the **mixture equations**, hold even if the densities in the mixture are *not Gaussian*.

Approximation of a Mixture

A mixture pdf can be approximated by a single Gaussian pdf with moments equal to those of the mixture, given by (1.4.16-6) and (1.4.16-8) — this is called **moment matching**.

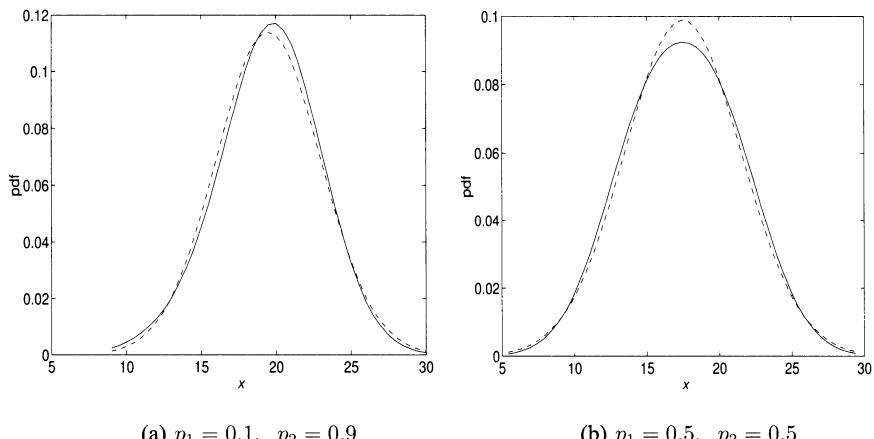


Figure 1.4.16-1: The exact Gaussian mixture pdf (solid line) and the corresponding *moment-matched Gaussian* (dashed line) for $\bar{x}_1 = 15, \bar{x}_2 = 20, \sigma_1^2 = \sigma_2^2 = 10$.

Figures 1.4.16-1 and 1.4.16-2 present comparisons between the exact pdf of Gaussian mixtures and the corresponding *moment-matched Gaussian* for several values of the parameters entering into the mixture.

In Fig. 1.4.16-1 the difference of the means is 5 while in Fig. 1.4.16-2 it is 10; in both cases the (common) standard deviation of the mixture components is $\sqrt{10}$.

The resulting “umbrella” Gaussian appears to be close to the exact pdf of the mixture as long as its components are not too far apart; namely, the distance between the means of the components is up to about *two standard deviations*.

This approximate condition is met for the cases depicted in Fig. 1.4.16-1, where the match appears good, and it is not met for the cases in Fig. 1.4.16-2, where the match does not appear good.

These results are used in obtaining recursive filtering algorithms for systems with hybrid uncertainties — *hybrid systems*. Systems falling into this category are those described by multiple models and/or with measurements of uncertain origin. (See also problem 1-4.)

1.4.17 Chi-Square Distributed Random Variables

If the n -dimensional random vector x is Gaussian, with mean \bar{x} and covariance P , then the (scalar) random variable defined by the quadratic form

$$q = (x - \bar{x})' P^{-1} (x - \bar{x}) \quad (1.4.17-1)$$

can be shown to be the sum of the squares of n independent zero-mean, unity-variance Gaussian random variables. Such a random variable is said to have a

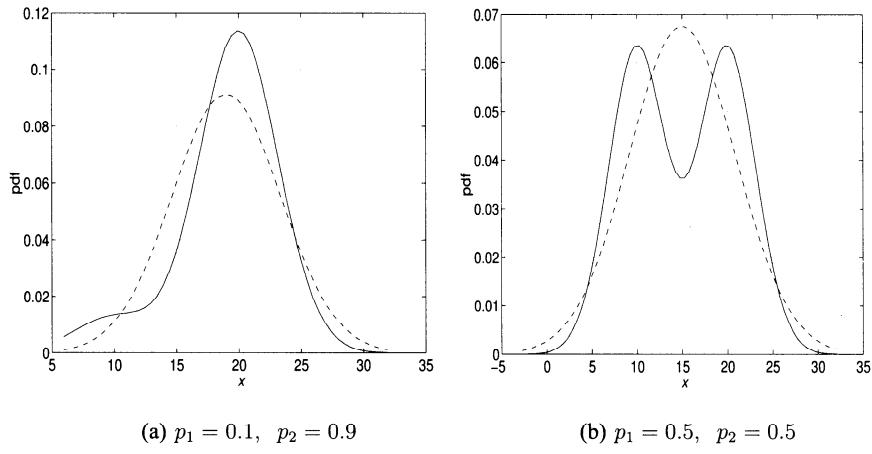


Figure 1.4.16-2: The exact Gaussian mixture pdf (solid line) and the corresponding moment-matched Gaussian (dashed line) for $\bar{x}_1 = 10$, $\bar{x}_2 = 20$, $\sigma_1^2 = \sigma_2^2 = 10$.

chi-square distribution with n degrees of freedom (the meaning of degrees of freedom is discussed in Subsection 3.6.2).

This can be seen as follows. Let

$$u \triangleq P^{-1/2}(x - \bar{x}) \quad (1.4.17-2)$$

Then u is Gaussian with

$$E[u] = 0 \quad (1.4.17-3)$$

$$E[uu'] = P^{-1/2} E[(x - \bar{x})(x - \bar{x})'] P^{-1/2} = P^{-1/2} P P^{-1/2} = I \quad (1.4.17-4)$$

where I is the identity matrix. Therefore, since the covariance matrix of u is diagonal, its components are uncorrelated and, in view of the fact that they are jointly Gaussian, they are also independent.

Thus

$$q = u'u = \sum_{i=1}^n u_i^2 \quad (1.4.17-5)$$

where, from (1.4.17-3) and (1.4.17-4), u_i are now **standard Gaussian random variables** (zero mean and unity variance)

$$u_i \sim \mathcal{N}(0, 1) \quad (1.4.17-6)$$

Therefore q is chi-square distributed with n degrees of freedom, which can be written as

$$q \sim \chi_n^2 \quad (1.4.17-7)$$

The mean and variance of the χ_n^2 random variable q are

$$E[q] = E\left[\sum_{i=1}^n u_i^2\right] = n \quad (1.4.17-8)$$

$$\begin{aligned}\text{var}(q) &= E\left[\sum_{i=1}^n (u_i^2 - 1)\right]^2 = \sum_{i=1}^n E[(u_i^2 - 1)^2] \\ &= \sum_{i=1}^n (E[u_i^4] - 2E[u_i^2] + 1) = \sum_{i=1}^n (3 - 2 + 1) = 2n \quad (1.4.17-9)\end{aligned}$$

where the fact that the cross-terms are zero mean has been used together with the expression of the fourth moment of a Gaussian variable, as given in (1.4.15-7).

The pdf of q — the chi-square density with n degrees of freedom — is

$$p(q) = \frac{1}{2^{\frac{n}{2}} \Gamma\left(\frac{n}{2}\right)} q^{\frac{n-2}{2}} e^{-\frac{q}{2}} \quad q \geq 0 \quad (1.4.17-10)$$

where Γ is the gamma function, with the following useful properties:

$$\Gamma\left(\frac{1}{2}\right) = \sqrt{\pi} \quad \Gamma(1) = 1 \quad \Gamma(m+1) = m\Gamma(m) \quad (1.4.17-11)$$

The square root of q has the **chi distribution**. The “chi-distributed” random variables with $n = 2$ and 3 degrees of freedom (and multiplied by an arbitrary constant) have the **Rayleigh distribution** and **Maxwell distribution**, respectively.

Given the *independent* random variables

$$q_1 \sim \chi_{n_1}^2 \quad q_2 \sim \chi_{n_2}^2 \quad (1.4.17-12)$$

then it can be easily shown that their sum

$$q_3 = q_1 + q_2 \quad (1.4.17-13)$$

is chi-square distributed with

$$n_3 = n_1 + n_2 \quad (1.4.17-14)$$

degrees of freedom. Symbolically, one can write

$$\chi_{n_1}^2 + \chi_{n_2}^2 = \chi_{n_1+n_2}^2 \quad (1.4.17-15)$$

The chi-square distribution is often used to check state estimators for “consistency” — that is, whether their actual errors are consistent with the variances calculated by the estimator.

1.4.18 Weighted Sum of Chi-Square Random Variables

Consider the following independent and identically distributed random variables

$$x_i \sim \chi_m^2 \quad i = 1, \dots, n \quad (1.4.18-1)$$

with m degrees of freedom chi-square distribution.

Then the **weighted sum of chi-square variables** with weights a_i ,

$$y_n \triangleq \sum_{i=1}^n a_i x_i \quad (1.4.18-2)$$

has a distribution which is not chi-square and is very complicated. Its mean is

$$E[y_n] = m \sum_{i=1}^n a_i \quad (1.4.18-3)$$

and, from (1.4.17-9), its variance is

$$\text{var}(y_n) = 2m \sum_{i=1}^n a_i^2 \quad (1.4.18-4)$$

The pdf of (1.4.18-2) can be approximated by **moment matching** — equating its first two moments to those of another random variable. In this case, the latter is chosen as a “scaled” chi-square with a number of degrees of freedom n' , to be determined together with the scaling factor.

To find a random variable whose moments are matched to those of (1.4.18-2), let

$$v \sim \chi_{n'}^2 \quad (1.4.18-5)$$

Then the mean and variance of a “scaled” version of v ,

$$w \triangleq cv \quad (1.4.18-6)$$

are cn' and $2c^2n'$, respectively. Equating these two moments to (1.4.18-3) and (1.4.18-4) yields the following equations for c and n' :

$$m \sum_{i=1}^n a_i = cn' \quad (1.4.18-7)$$

$$2m \sum_{i=1}^n a_i^2 = 2c^2n' \quad (1.4.18-8)$$

The solution of these equations yields the scaling factor as

$$c = \frac{\sum_{i=1}^n a_i^2}{\sum_{i=1}^n a_i} \quad (1.4.18-9)$$

and the number of degrees of freedom of v as

$$n' = \frac{m (\sum_{i=1}^n a_i)^2}{\sum_{i=1}^n a_i^2} \quad (1.4.18-10)$$

Therefore, the distribution of y_n , defined in (1.4.18-2) is approximately

$$y_n \sim \frac{\sum_{i=1}^n a_i^2}{\sum_{i=1}^n a_i} \chi_{n'}^2 \quad (1.4.18-11)$$

Note that the χ_n^2 density was defined only for integer n ; a noninteger n makes it into a case of the gamma density. In practice, one can interpolate from the chi-square tables.

Fading Memory Average

Consider the case where the weights in (1.4.18-2) are exponential, that is,

$$a_i = \alpha^{n-i} \quad (1.4.18-12)$$

with $0 < \alpha < 1$. Then the weighted sum

$$z_n \triangleq \sum_{i=1}^n \alpha^{n-i} x_i \quad (1.4.18-13)$$

is called the *fading memory sum* of the variables x_i .

Then, for $n \gg 1$ one obtains

$$z_n \sim \frac{1}{1 + \alpha} \chi_{n'}^2 \quad (1.4.18-14)$$

and

$$n' = m \frac{1 + \alpha}{1 - \alpha} \quad (1.4.18-15)$$

When (1.4.18-13) is normalized by dividing it with the sum of the coefficients, it becomes the *fading memory average* or *exponentially discounted average*.

These results are used in real-time monitoring of state estimator performance.

1.4.19 Random Processes

A scalar random variable is a (real) number x determined by the outcome ω of a random experiment

$$x = x(\omega) \quad (1.4.19-1)$$

A (scalar) *random process* or a *stochastic process* is a *function of time* determined by the outcome of a random experiment

$$x(t) = x(t, \omega) \quad (1.4.19-2)$$

This is a family or **ensemble** of functions of time, in general different for each outcome ω .

The **mean** or **ensemble average** of the random process is

$$\bar{x}(t) = E[x(t)] = \int_{-\infty}^{\infty} \xi p_{x(t)}(\xi) d\xi \quad (1.4.19-3)$$

while its **autocorrelation** is defined for a real-valued (scalar) process as

$$R(t_1, t_2) \triangleq E[x(t_1)x(t_2)] = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} \xi \eta p_{x(t_1), x(t_2)}(\xi, \eta) d\xi d\eta \quad (1.4.19-4)$$

The **autocovariance** of this random process is

$$\begin{aligned} V(t_1, t_2) &\triangleq E[[x(t_1) - \bar{x}(t_1)][x(t_2) - \bar{x}(t_2)]] \\ &= R(t_1, t_2) - \bar{x}(t_1)\bar{x}(t_2) \end{aligned} \quad (1.4.19-5)$$

Note that the autocorrelation of a random process is an (unnormalized) *non-central moment* — the mean is not subtracted — while the correlation of two random variables is the normalized joint *central moment* (1.4.6-6). If the process is zero mean, the distinction between central and noncentral moments disappears.

For a **vector-valued random process**, (1.4.19-4) and (1.4.19-5) contain the corresponding outer products.

Stationarity

A random process whose mean is time invariant and whose autocorrelation is of the form

$$R(t_1, t_2) = R(t_1 - t_2) = R(\tau) \quad (1.4.19-6)$$

where

$$\tau = t_1 - t_2 \quad (1.4.19-7)$$

that is, its first two moments are invariant with respect to a shift of the time axis, is called **wide sense stationary** or, less rigorously, **stationary**.

As can be seen from (1.4.19-4), for a real-valued process the autocorrelation is symmetric in its two time arguments

$$R(t_1, t_2) = R(t_2, t_1) \quad (1.4.19-8)$$

Thus

$$R(\tau) = R(-\tau) \quad (1.4.19-9)$$

The **power spectrum** or **power spectral density** of a *stationary* random process is the Fourier transform (denoted by \mathcal{F}) of its autocorrelation

$$S(\omega) = \mathcal{F}\{R(\tau)\} = \int_{-\infty}^{\infty} e^{-j\omega\tau} R(\tau) d\tau \quad (1.4.19-10)$$

where ω denotes the (angular) frequency.

Strict stationarity requires that all pdfs (rather than only moments up to second order) be invariant with respect to time shift. In practice, wide-sense stationarity is about all one can assume (and hope for).

Example

Consider the scalar dynamic system

$$\dot{x}(t) = -ax(t) + n(t) \quad (1.4.19-11)$$

with $a > 0$, i.e., a stable system. The input $n(t)$ is a zero-mean white noise with autocorrelation (same as its autocovariance, since it is zero-mean)

$$E[n(t_1)n(t_2)] = S_0\delta(t_1 - t_2) \quad (1.4.19-12)$$

Assuming that the system in steady state (i.e., the initial condition died out), we obtain

$$x(t) = \int_{-\infty}^t h(t, \tau)n(\tau) d\tau \quad (1.4.19-13)$$

where the impulse response of the system considered is

$$h(t, \tau) = e^{-a(t-\tau)} \quad (1.4.19-14)$$

The mean of the output $x(t)$ is

$$E[x(t)] = E\left[\int_{-\infty}^t h(t, \tau)n(\tau) d\tau\right] = \int_{-\infty}^t h(t, \tau)E[n(\tau)] d\tau = 0 \quad (1.4.19-15)$$

The autocorrelation of $x(t)$ is

$$E[x(t_1)x(t_2)] = E\left[\int_{-\infty}^{t_1} h(t_1, \tau_1)n(\tau_1) d\tau_1 \int_{-\infty}^{t_2} h(t_2, \tau_2)n(\tau_2) d\tau_2\right] \quad (1.4.19-16)$$

Assuming $t_1 < t_2$, the above can be written as

$$\begin{aligned} E[x(t_1)x(t_2)] &= E\left\{\int_{-\infty}^{t_1} h(t_1, \tau_1)n(\tau_1) d\tau_1 \cdot \right. \\ &\quad \left. \left[\int_{-\infty}^{t_1} h(t_2, \tau_2)n(\tau_2) d\tau_2 + \int_{t_1}^{t_2} h(t_2, \tau_2)n(\tau_2) d\tau_2\right]\right\} \end{aligned} \quad (1.4.19-17)$$

In view of the fact that the first and last integral above are over nonoverlapping intervals, the expected value of their product will be zero due to the whiteness of $n(t)$, whose autocorrelation is zero at any two different times.

Thus one can write

$$\begin{aligned} E[x(t_1)x(t_2)] &= E\left\{\int_{-\infty}^{t_1} e^{-a(t_1-\tau_1)}n(\tau_1) d\tau_1 \int_{-\infty}^{t_1} e^{-a(t_2-\tau_2)}n(\tau_2) d\tau_2\right\} \\ &= e^{a(t_1-t_2)}E\left\{\int_{-\infty}^{t_1} e^{-a(t_1-\tau_1)}n(\tau_1) d\tau_1 \int_{-\infty}^{t_1} e^{-a(t_1-\tau_2)}n(\tau_2) d\tau_2\right\} \\ &= e^{a(t_1-t_2)}E[x(t_1)^2] \end{aligned} \quad (1.4.19-18)$$

If $t_1 > t_2$ the result would be the same except that they are switched. Therefore,

$$E[x(t_1)x(t_2)] = e^{-a|t_1-t_2|}E[x(t_1)^2] \quad (1.4.19-19)$$

i.e., the autocorrelation decreases exponentially from its peak, which is the (steady-state) mean square value of $x(t)$ — its instantaneous power.

The mean square value of $x(t)$ is

$$\begin{aligned} E[x(t)^2] &= E \left\{ \int_{-\infty}^t e^{-a(t-\tau_1)} n(\tau_1) d\tau_1 \int_{-\infty}^t e^{-a(t-\tau_2)} n(\tau_2) d\tau_2 \right\} \\ &= \int_{-\infty}^t \int_{-\infty}^t e^{-a(2t-\tau_1-\tau_2)} E[n(\tau_1)n(\tau_2)] d\tau_1 d\tau_2 \\ &= e^{-2at} \int_{-\infty}^t \int_{-\infty}^t e^{a(\tau_1+\tau_2)} S_0 \delta(\tau_1 - \tau_2) d\tau_1 d\tau_2 \\ &= e^{-2at} S_0 \int_{-\infty}^t e^{2a\tau_1} d\tau_1 = e^{-2at} \frac{S_0}{2a} [e^{2at} - 0] = \frac{S_0}{2a} \quad (1.4.19-20) \end{aligned}$$

Thus

$$R_{xx}(\tau) = E[x(t)x(t+\tau)] = \frac{S_0}{2a} e^{-a|\tau|} \quad (1.4.19-21)$$

The Fourier transform of the above, using the unit step function $1(\cdot)$ to rewrite the two-sided exponential, yields the spectrum of x as

$$\begin{aligned} S_{xx}(\omega) &= \mathcal{F}\{R_{xx}(\tau)\} = \frac{S_0}{2a} \mathcal{F}\{e^{a\tau} 1(-\tau) + e^{-a\tau} 1(\tau)\} \\ &= \frac{S_0}{2a} \left[\frac{1}{a - j\omega} + \frac{1}{a + j\omega} \right] = S_0 \frac{1}{a^2 + \omega^2} \quad (1.4.19-22) \end{aligned}$$

Since this stable system acts as a low-pass filter, its effect on the spectrum of the input (which, being a white noise, has a flat spectrum) is to attenuate the high frequencies.

Ergodicity

If a stationary random process is **ergodic**, then **time averages** (of some functions of the process) are equal to the corresponding expected values — **ensemble averages** — given by (1.4.19-3), that is,

$$\lim_{T \rightarrow \infty} \frac{1}{2T} \int_{-T}^T x(t) dt = \bar{x} \quad (1.4.19-23)$$

Time averages are used in real-time performance monitoring of state estimation filters.

White Noise

A (not necessarily stationary) random process whose autocovariance is zero for any two different times is called **white noise**. In this case

$$V(t_1, t_2) = \sigma^2(t_1) \delta(t_1 - t_2) = S_0(t_1) \delta(t_1 - t_2) \quad (1.4.19-24)$$

where $\sigma^2(t_1)$ is its “instantaneous variance” and $\delta(\cdot)$ is the **Dirac (impulse) delta function**. The above property is the **wide-sense whiteness** as opposed to the **strict-sense whiteness**, which is defined by *independence* rather than *uncorrelatedness* as in (1.4.19-24).

A **stationary zero-mean white** process has the autocorrelation

$$R(\tau) = E[x(t + \tau)x(t)] = S_0\delta(\tau) \quad (1.4.19-25)$$

The impulse function autocorrelation (1.4.19-25) leads to a power spectral density S_0 constant across the frequency spectrum.

It is convenient in some applications to consider a **nonstationary zero-mean white** process $x(t)$. The autocorrelation of such a process is

$$E[x(t_1)x(t_2)] = S_0(t_1)\delta(t_1 - t_2) \quad (1.4.19-26)$$

where $S_0(t)$ is, with some abuse of language, the “time-varying spectral density” or the “instantaneous variance.” The term **instantaneous variance** is in fact *incorrect* in this case since the variance is actually *infinite*.¹¹

Note on Whiteness

In the sequel, when a result requires whiteness, it usually requires strict-sense whiteness. In practice, however, one has only moments (typically up to second order) and thus only wide-sense whiteness can be assumed. Then the result is of practical value only when the strict sense whiteness is replaced by wide-sense whiteness, although in this case the result is only approximately true in theory.

1.4.20 Random Walk and the Wiener Process

The **Wiener random process** (or Wiener-Levy or Brownian motion) is a limiting form of the **random walk**: the sum of independent steps of size $s \rightarrow 0$, equiprobable in each direction, taken at intervals $\Delta \rightarrow 0$ such that

$$\frac{s}{\sqrt{\Delta}} \rightarrow \sqrt{\alpha} \quad (1.4.20-1)$$

where α is a constant. This yields a stochastic process $\mathbf{w}(t)$ with the following pdf [assuming $\mathbf{w}(0) = 0$],

$$p[\mathbf{w}(t)] = \mathcal{N}[\mathbf{w}(t); 0, \alpha t] \quad (1.4.20-2)$$

that is, normal, zero mean, and with variance αt .

¹¹Also, the physical dimension of σ^2 is the square of the dimension of x , *multiplied by time*, since δ has as dimension the inverse of time.

Note that the Wiener process is *nonstationary*. It relates to the zero-mean white noise, denoted here as $n(t)$, as follows

$$\mathbf{w}(t) = \int_0^t n(\tau) d\tau \quad (1.4.20-3)$$

where

$$E[n(t_1)n(t_2)] = \alpha\delta(t_1 - t_2) \quad (1.4.20-4)$$

Another way of writing (1.4.20-3) is

$$d\mathbf{w}(t) = n(t) dt \quad (1.4.20-5)$$

which shows that the Wiener process is an *independent increment process*. Formally, the white noise is the derivative of the Wiener process; however, this is not rigorous, since the Wiener process is *nowhere differentiable* — its derivative has infinite variance.

The autocorrelation of the Wiener process is

$$E[\mathbf{w}(t_1)\mathbf{w}(t_2)] = \alpha \min(t_1, t_2) \quad (1.4.20-6)$$

The white noise and the Wiener process are used to model unknown inputs (maneuvers) in state estimation/tracking.

Note

Since the Wiener process is not stationary, it does not have a power spectral density.

1.4.21 Markov Processes

Markov processes are defined by the following *Markov property*

$$p[x(t)|x(\tau), \tau \leq t_1] = p[x(t)|x(t_1)] \quad \forall t > t_1 \quad (1.4.21-1)$$

that is, the past up to any t_1 is *fully characterized* by the value of the process at t_1 .

An equivalent statement to the above is:

"The future is independent of the past if the present is known."

The Wiener process is Markov. This follows from the fact that it is the integral of white noise

$$\mathbf{w}(t) = \mathbf{w}(t_1) + \int_{t_1}^t n(\tau) d\tau \quad (1.4.21-2)$$

and $n(\tau), \tau \in [t_1, t]$ is independent of $\mathbf{w}(t_1)$.

Furthermore, the state $x(t)$ of a (possibly time-varying) dynamic system driven by white noise $n(t)$,

$$\dot{x}(t) = f[t, x(t), n(t)] \quad (1.4.21-3)$$

is a Markov process. In general, both $x(t)$ and $n(t)$ are vector-valued random processes.

Markov Processes with Rational Spectra

Given a *linear time-invariant dynamic system excited by a stationary white noise $n(t)$ with mean zero*

$$\dot{x}(t) = Ax(t) + Bn(t) \quad (1.4.21-4)$$

its state is (in steady state — if the system is stable) a *stationary Markov process* with spectrum (power spectral density)

$$S(\omega) = H(j\omega)QH(j\omega)^* \quad (1.4.21-5)$$

In the above the asterisk denotes complex conjugate transpose and

$$H(j\omega) = (j\omega I - A)^{-1}B \quad (1.4.21-6)$$

is the *transfer function matrix* (from the input to the state) of the system (1.4.21-4) and the autocorrelation of the input (in general, a matrix) is

$$R_n(t_1, t_2) = E[n(t_1)n(t_2)'] = Q\delta(t_1 - t_2) \quad (1.4.21-7)$$

The matrix Q is sometimes called (nonrigorously) the covariance of $n(t)$; actually it is its *power spectral density*.

Note that (1.4.21-5) is a *rational spectrum* — a ratio of polynomials.

Conversely, one has the following result:

Every Markov process with a rational spectrum can be represented as a linear time-invariant system excited by white noise.

The models used in state estimation are Markov processes — linear or nonlinear systems driven by white noise. These systems can be time-invariant or time-varying and the noise can be stationary or nonstationary. If a time-invariant system is driven by a stationary noise that is not white, but has a rational spectrum, then it has to undergo *prewhitening*: It is the output of a subsystem driven by white noise. The transfer function of this subsystem is obtained by *spectral factorization* — decomposition of the output spectrum in the form (1.4.21-5). For more details on spectral factorization, see, e.g., [Poor88].

Prewhitenning

The prewhitening is illustrated in Figs. 1.4.21-1 and 1.4.21-2. In Fig. 1.4.21-1 the system \mathcal{S} with state $x(t)$ is driven by the *autocorrelated noise* (also called *colored noise*) $n(t)$. Since $n(t)$ is not white, $x(t)$ is not a Markov process.

Assume that $n(t)$ can be represented as the output of a system \mathcal{S}_0 with white noise input $v(t)$. With the state of the *prewhitening system or shaping filter \mathcal{S}_0* denoted as $x_0(t)$, the augmented state of the composite system shown in Fig. 1.4.21-2,

$$y(t) \triangleq \begin{bmatrix} x(t) \\ x_0(t) \end{bmatrix} \quad (1.4.21-8)$$

is then a Markov process.

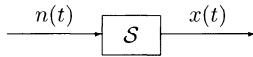


Figure 1.4.21-1: A system driven by autocorrelated noise.

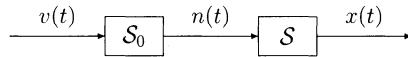


Figure 1.4.21-2: The same system augmented to be driven by white noise.

Example of Spectral Factorization

Consider the random process $x(t)$ with the spectrum

$$S_{xx}(\omega) = S_0 \frac{1}{a^2 + \omega^2} \quad (1.4.21-9)$$

This can be factorized as follows. Consider the random process $x(t)$ with the spectrum

$$S_{xx}(\omega) = \frac{1}{a + j\omega} S_0 \frac{1}{a - j\omega} \quad (1.4.21-10)$$

In the above, one can identify the *causal* transfer function

$$H(\omega) = \frac{1}{a + j\omega} \quad (1.4.21-11)$$

as the transfer function of the system

$$\dot{x}(t) = -ax(t) + n(t) \quad (1.4.21-12)$$

driven by the white noise $n(t)$.

If $n(t)$ is zero-mean with spectrum S_0 , then the above system is the *prewhitening system* for the autocorrelated process $x(t)$.

Note that in the course of the factorization, the roles of H and its complex conjugate H^* can be switched. It is, however, necessary to choose H such that it represents a stable and causal system. Namely, the impulse response, given by the inverse Fourier transform, denoted by \mathcal{F}^{-1} , of the transfer function, i.e.,

$$\mathcal{F}^{-1} \left\{ \frac{1}{a + j\omega} \right\} = e^{-at} 1(t) \quad (1.4.21-13)$$

is *causal*. The impulse response corresponding to the other factor, namely,

$$\mathcal{F}^{-1} \left\{ \frac{1}{a - j\omega} \right\} = e^{at} 1(-t) \quad (1.4.21-14)$$

is *noncausal* (cannot be implemented by a physical system) and thus cannot serve as H .

1.4.22 Random Sequences, Markov Sequences and Markov Chains

A *random sequence*, or a *discrete-time stochastic process*, is a time-indexed sequence of random variables

$$X^k = \{x(j)\}_{j=1}^k \quad k = 1, 2, \dots \quad (1.4.22-1)$$

Similarly to the continuous-time definition of the Markov property, a random sequence is Markov if

$$p[x(k)|X^j] = p[x(k)|x(j)] \quad \forall k > j \quad (1.4.22-2)$$

The (real-valued) *zero-mean sequence* $v(j)$, $j = 1, \dots$, is a *discrete-time white noise* (a *white sequence*) if

$$E[v(k)v(j)] = q(k)\delta_{kj} \quad (1.4.22-3)$$

where the *Kronecker delta function*

$$\delta_{kj} = \begin{cases} 1 & \text{if } k = j \\ 0 & \text{if } k \neq j \end{cases} \quad (1.4.22-4)$$

is used and $q(k)$ denotes its variance. If $q(k) = q$, that is, the variance is time-invariant, then this is a *stationary white sequence*.

The sequence with property (1.4.22-3) is actually only uncorrelated, or weakly independent — strong independence requires that (1.4.7-3) be satisfied. However, in practice one does not have pdfs but only moments up to second order, and thus the usual assumption is the *weak independence (uncorrelatedness)* indicated in (1.4.22-3).

The state $x(k)$ of a dynamic system excited by white noise $v(k)$

$$x(k+1) = f[k, x(k), v(k)] \quad (1.4.22-5)$$

is a *discrete-time Markov process* or *Markov sequence*. In general, both $x(k)$ and $v(k)$ are vector-valued.

The state of a linear dynamic system excited by white Gaussian noise

$$x(k+1) = F(k)x(k) + v(k) \quad (1.4.22-6)$$

is a *Gauss-Markov sequence*.

The reasons for this are as follows: Because of the linearity, $x(k)$ is Gaussian (assuming the initial condition is Gaussian); and because of the whiteness of the driving (“process”) noise, it is Markov.

A special case of (1.4.22-6), for a scalar x , is

$$x(k+1) = x(k) + v(k) \quad (1.4.22-7)$$

in which case x becomes the *integral (sum)* of the white noise sequence terms, and is called a *discrete-time Wiener process*.

Markov Chains

A **Markov chain** is a special case of a Markov sequence, in which the state space is discrete and finite:

$$x(k) \in \{x_i, i = 1, \dots, n\} \quad (1.4.22-8)$$

Its characterization is given in full by the **transition (jump) probabilities**

$$P\{x(k) = x_j | x(k-1) = x_i\} \triangleq \pi_{ij} \quad i, j = 1, \dots, n \quad (1.4.22-9)$$

and the initial probabilities.

Define the vector

$$\mu(k) \triangleq [\mu_1(k), \dots, \mu_n(k)]' \quad (1.4.22-10)$$

where the components are the probabilities of the chain being in state i

$$\mu_i(k) \triangleq P\{x(k) = x_i\} \quad (1.4.22-11)$$

This vector describes the pmf of the state of the chain.

The evolution in time of (1.4.22-11) is then given by

$$\mu_i(k+1) = \sum_{j=1}^n \pi_{ji} \mu_j(k) \quad i = 1, \dots, n \quad (1.4.22-12)$$

It can be easily shown that the above can be written in vector form with notation (1.4.22-10) as

$$\mu(k+1) = \Pi' \mu(k) \quad (1.4.22-13)$$

where

$$\Pi = [\pi_{ij}] \quad (1.4.22-14)$$

is the **transition matrix of the Markov chain**.

1.4.23 The Law of Large Numbers and the Central Limit Theorem

The Law of Large Numbers

The **law of large numbers (LLN)** states loosely that the sum of a large number of random variables tends, under some fairly nonrestrictive conditions, to its expected value. One of the versions of the LLN is the following.

Given a stationary sequence of random variables $x_i, i = 1, \dots$, with means

$$E[x_i] = \bar{x} \quad (1.4.23-1)$$

and covariances

$$E[(x_i - \bar{x}_i)(x_j - \bar{x}_j)] = \sigma^2 \rho(i-j) \quad (1.4.23-2)$$

where the correlation coefficients are such that

$$\lim_{|i-j| \rightarrow \infty} \rho(i-j) = 0 \quad (1.4.23-3)$$

that is, this sequence is *asymptotically uncorrelated*.

If the correlation coefficients in (1.4.23-3) tend to zero “sufficiently fast” (e.g., exponential decay), then the *sample average*

$$y_n \triangleq \frac{1}{n} \sum_{i=1}^n x_i \quad (1.4.23-4)$$

converges, as $n \rightarrow \infty$, in the mean square sense to its expected value

$$\bar{y}_n = \frac{1}{n} \sum_{i=1}^n \bar{x} = \bar{x} \quad (1.4.23-5)$$

In other words, the variance of y_n in (1.4.23-4) tends to zero as n increases.

If the stronger condition of uncorrelatedness holds, then the convergence of y_n given by (1.4.23-4) to \bar{y}_n given by (1.4.23-5) is quite obvious — it can be easily shown that in this case the mean square value of their difference is σ^2/n .

The Central Limit Theorem

The *central limit theorem (CLT)* states that if the sequence x_i , $i = 1, \dots$, consists of *independent* random variables, then under some reasonably mild conditions the pdf of the sum

$$z_n = \frac{1}{\sqrt{n}} \sum_{i=1}^n x_i \quad (1.4.23-6)$$

will tend to a Gaussian pdf as $n \rightarrow \infty$.

If these random variables are *independent and identically distributed (i.i.d.)* zero-mean and with variance σ^2 , then the z_n tends to $\mathcal{N}(0, \sigma^2)$.

Since y_n defined in (1.4.23-4) is a scaled version of z_n , the central limit theorem holds for y_n as well. Thus, if the random variables x_i are i.i.d., then for large n ,

$$y_n \sim \mathcal{N}(\bar{y}_n, \sigma_{y_n}^2) \quad (1.4.23-7)$$

where

$$\bar{y}_n = \bar{x} \quad (1.4.23-8)$$

$$\sigma_{y_n}^2 = \frac{1}{n^2} \sum_{i=1}^n \sigma^2 = \frac{\sigma^2}{n} \quad (1.4.23-9)$$

are the mean and variance of y_n .

The independence requirement can be weakened — there are versions of the central limit theorem that allow dependence, as long as the random variables become “asymptotically independent,” in a manner somewhat similar to the asymptotic uncorrelatedness property (1.4.23-3).

Note on the CLT

The CLT has a very important role in characterizing many real-world sources of uncertainty: It is used as the justification/excuse to make the omnipresent Gaussian assumption. For example, the thermal noise in electronic devices, as the sum of many “small contributions,” is indeed close to Gaussian.

1.5 BRIEF REVIEW OF STATISTICS

1.5.1 Hypothesis Testing

Consider the *hypothesis testing* problem between the following “simple” hypotheses — each is defined by a point (a certain value of a parameter).

The *null hypothesis* is

$$H_0 : \theta = \theta_0 \quad (1.5.1-1)$$

while the *alternate hypothesis* is

$$H_1 : \theta = \theta_1 \quad (1.5.1-2)$$

where θ is a certain parameter whose value is equal to a certain value θ_0 (θ_1) under H_0 (H_1).

The *type I error* probability is defined as

$$P_{e_I} \triangleq P\{\text{“}H_1\text{”}|H_0\} \triangleq P\{\text{accept } H_1|H_0 \text{ true}\} \quad (1.5.1-3)$$

while the *type II error* probability is

$$P_{e_{II}} \triangleq P\{\text{“}H_0\text{”}|H_1\} \triangleq P\{\text{accept } H_0|H_1 \text{ true}\} \quad (1.5.1-4)$$

In signal detection, if H_0 stands for “signal equal to zero” (i.e., absent) and H_1 stands for “signal present,” then the type I error is a *false alarm* while the type II error is a *miss*.

The *power of the test* is

$$\pi \triangleq P\{\text{“}H_1\text{”}|H_1\} = 1 - P_{e_{II}} \quad (1.5.1-5)$$

and it measures the test’s capability to discern H_1 when it is true — the *detection probability*.

The decision as to which hypothesis to accept is made based on a set of observations, Z , whose pdfs conditioned on H_0 and H_1 are known.

According to the *Neyman-Pearson Lemma*, the optimal decision, in the sense of *minimizing the probability of type II error* (or maximizing the power of the test), *subject to a given (maximum) probability of type I error* is as follows.

The test, based on the *likelihood ratio*, is

$$\Lambda(H_1, H_0) = \frac{p(Z|H_1)}{p(Z|H_0)} \stackrel{\substack{"H_1" \\ "H_0" \\ \geq}}{\gtrless} \Lambda_0 \quad (1.5.1-6)$$

that is, “accept H_1 ” if Λ exceeds the threshold Λ_0 and “accept H_0 ” if Λ is below this threshold. The threshold Λ_0 is such that

$$P\{\Lambda(H_1, H_0) > \Lambda_0 | H_0\} = P_{e_I} \quad (1.5.1-7)$$

The pdf of the observations, $p(Z|H_j)$, is called the *likelihood function* of H_j .

Example

Consider the test between the hypotheses (1.5.1-1) and (1.5.1-2) with $\theta_0 = 0$ and $\theta_1 > 0$, based on the single observation

$$z \sim \mathcal{N}(\theta, \sigma^2) \quad (1.5.1-8)$$

that is, we are testing whether the mean of z is zero or positive.

The likelihood ratio is

$$\Lambda(H_1, H_0) = \frac{p(z|H_1)}{p(z|H_0)} = e^{-\frac{(z-\theta_1)^2-z^2}{2\sigma^2}} = e^{\frac{2z\theta_1-\theta_1^2}{2\sigma^2}} \quad (1.5.1-9)$$

Instead of comparing (1.5.1-9) to a threshold, it is convenient to take its logarithm, incorporate the various constants into the threshold, and then determine it.

Since $\theta_1 > 0$, comparing (1.5.1-9) to a threshold is equivalent to

$$z \stackrel{\substack{"H_1" \\ "H_0" \\ \geq}}{\gtrless} \lambda_0 \quad (1.5.1-10)$$

where λ_0 is the new threshold to be determined.

In view of (1.5.1-8), the pdf of the observation under H_0 is

$$p(z|H_0) = \mathcal{N}(z; 0, \sigma^2) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{z^2}{2\sigma^2}} \quad (1.5.1-11)$$

The threshold λ_0 for *false alarm probability* $P_{FA} = \alpha$ follows from

$$P_{FA} \triangleq P_{e_I} = P\{z > \lambda_0 | H_0\} = \int_{\lambda_0}^{\infty} \mathcal{N}(z; 0, \sigma^2) dz = \alpha \quad (1.5.1-12)$$

The *tail mass* or *tail probability* α from (1.5.1-12) is shaded in Fig. 1.5.1-1.

Under H_0 there is only a small probability α that z falls in the shaded tail region of the likelihood function of H_0 ; if it does fall there, then it is deemed

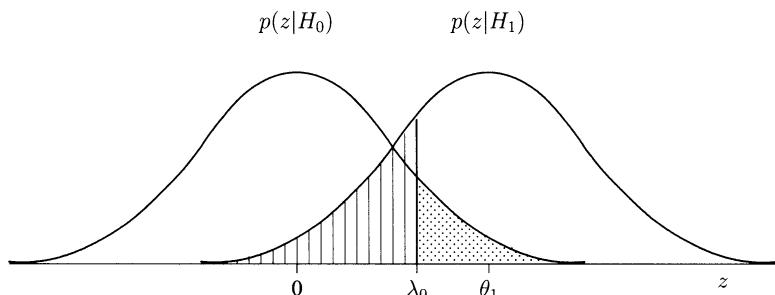


Figure 1.5.1-1: Test between two simple hypotheses.

unlikely that H_0 is true, and it is rejected. In this case it is said that the *level of significance* of H_0 is low: α or less.¹² The latter holds if $z = z_1 > \lambda_0$ because such an observation belongs to an even smaller tail mass.

The type II error probability — the *miss probability* — of the test (1.5.1-10) is

$$P_{e_{II}} \triangleq P\{z < \lambda_0 | H_1\} = \int_{-\infty}^{\lambda_0} \mathcal{N}(z; \theta_1, \sigma^2) dz \quad (1.5.1-13)$$

This corresponds to the striped area of Fig. 1.5.1-1.

The power of this test — the *probability of detection* — also denoted as P_D is

$$P_D \triangleq \pi = P\{z > \lambda_0 | H_1\} = \int_{\lambda_0}^{\infty} \mathcal{N}(z; \theta_1, \sigma^2) dz = 1 - P_{e_{II}} \quad (1.5.1-14)$$

and, obviously, it depends on θ_1 — the farther it is from 0, the more powerful the test will be.

Example

For $\alpha = 5\%$, the threshold is obtained (see Subsection 1.5.4) based on the one-sided tail of the Gaussian pdf as $\lambda_0 = 1.64\sigma$. Table 1.5.1-1 shows the power of this test as a function of the separation between the two means (in units of σ , the common standard deviation for both hypotheses).

1.5.2 Confidence Regions and Significance

Assume one desires to test the null hypothesis

$$H_0 : \theta = 0 \quad (1.5.2-1)$$

¹²Equivalently, one can say that the *confidence level* for H_1 is high ($1 - \alpha$ or more).

Table 1.5.1-1: Power of the test between simple hypotheses as a function of their separation (threshold set for $\alpha = 5\%$).

θ_1/σ	2	3	4	5
Power π	0.641	0.913	0.991	0.9996

versus the **one-sided alternative hypothesis**

$$H'_1 : \theta > 0 \quad (1.5.2-2)$$

based upon a set of observations

$$Z = \{z_i, i = 1, \dots, n\} \quad (1.5.2-3)$$

where

$$z_i = \theta + w_i \quad (1.5.2-4)$$

and the “noises” w_i are independent and identically distributed:

$$w_i \sim \mathcal{N}(0, \sigma^2) \quad (1.5.2-5)$$

Note that H'_1 is a **composite hypothesis** — it is defined by more than one point (an interval in this case).

The test is subject to

$$P_{e_I} = \alpha \quad (1.5.2-6)$$

The likelihood ratio for hypothesis H'_1 versus H_0 is

$$\Lambda(H'_1, H_0) = \frac{p(Z|\theta)}{p(Z|\theta=0)} = e^{\frac{1}{2\sigma^2} \sum_{i=1}^n [z_i^2 - (z_i - \theta)^2]} = e^{\frac{1}{2\sigma^2} \sum_{i=1}^n (2z_i\theta - \theta^2)} \quad (1.5.2-7)$$

The comparison of the above to a threshold is equivalent to comparing the sample mean of the observations to another threshold:

$$\bar{z} \triangleq \frac{1}{n} \sum_{i=1}^n z_i \underset{\text{"H}_0\text{"}}{\underset{\text{"H}'_1\text{"}}{\gtrless}} \lambda_1 \quad (1.5.2-8)$$

Equation (1.5.2-8) follows by taking the logarithm of (1.5.2-7) and lumping θ and σ into the (yet undetermined) threshold; the division by n is done for the convenience of normalization — this yields \bar{z} in the same range regardless of the number of samples n .

For this problem \bar{z} is the **test statistic** — the function of the observations used in the test.

The threshold λ_1 is obtained by noting that

$$p(\bar{z}|H_0) = \mathcal{N}(\bar{z}; 0, \frac{\sigma^2}{n}) \quad (1.5.2-9)$$

and putting the condition

$$P\{\text{``}H'_1\text{''}|H_0\} = P\{\bar{z} > \lambda_1|H_0\} = 1 - \int_{-\infty}^{\lambda_1} \mathcal{N}(\bar{z}; 0, \frac{\sigma^2}{n}) d\bar{z} = \alpha \quad (1.5.2-10)$$

If the one-sided alternate hypothesis (1.5.2-2) is replaced by the ***two-sided alternate hypothesis***, denoted as H_1 ,

$$H_1 : \theta \neq 0 \quad (1.5.2-11)$$

then it can be shown that (1.5.2-8) is to be replaced by

$$|\bar{z}| \stackrel{\text{``}H_1\text{''}}{\underset{\text{``}H_0\text{''}}{\gtrless}} \lambda \quad (1.5.2-12)$$

The threshold λ is obtained by putting the condition

$$P\{\text{``}H_1\text{''}|H_0\} = P\{|\bar{z}| > \lambda|H_0\} = 1 - \int_{-\lambda}^{\lambda} \mathcal{N}(\bar{z}; 0, \frac{\sigma^2}{n}) d\bar{z} = \alpha \quad (1.5.2-13)$$

In other words, λ is such that the ***acceptance region*** for H_0 , which is the interval $[-\lambda, \lambda]$, contains $1 - \alpha$ probability mass for the pdf (1.5.2-9).

For $\alpha = 0.05$, one obtains from tables of the normal distribution (see Sub-section 1.5.4)

$$\lambda = 1.96 \frac{\sigma}{\sqrt{n}} \quad (1.5.2-14)$$

The power of the test depends on the specific value of θ under the alternate hypothesis H_1 .

The sample mean \bar{z} is also the ***maximum likelihood estimate*** (and the ***least squares estimate***) of the unknown parameter θ , i.e.,

$$\bar{z} = \arg \max_{\theta} p(Z|\theta) = \arg \min_{\theta} \sum_{i=1}^n (z_i - \theta)^2 \quad (1.5.2-15)$$

since $p(Z|\theta)$ is the likelihood function of the parameter θ .

Figure 1.5.2-1 illustrates the acceptance region for H_0 — the ***two-sided 95% probability region*** for \bar{z} under hypothesis H_0 — the interval $[-\lambda, \lambda]$. The shaded areas represent the tails on the two sides of $p(\bar{z}|H_0)$, whose total probability mass is $\alpha = 0.05$; this is the region of rejection of H_0 (acceptance of H_1).

Since the difference between the sample mean \bar{z} and the true mean θ (whatever its value) is

$$\bar{z} - \theta \sim \mathcal{N}(0, \frac{\sigma^2}{n}) \quad (1.5.2-16)$$

one can say that the true mean lies within the interval $[\bar{z} - \lambda, \bar{z} + \lambda]$ with “confidence” $1 - \alpha$ (since θ is not a random variable, no probabilistic statement about it can be made). In view of this, such an interval is called the ***confidence region*** for θ .

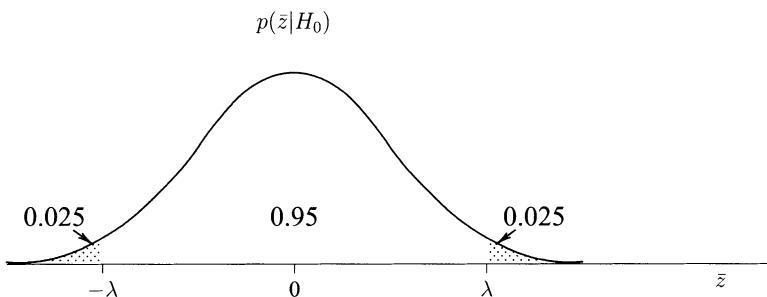


Figure 1.5.2-1: The test between hypotheses H_0 and H_1 .

Another way of interpreting test (1.5.2-12) is the following: If the null hypothesis value $\theta = 0$ falls within the confidence region $[\bar{z} - \lambda, \bar{z} + \lambda]$, then H_0 is accepted. This is equivalent to \bar{z} falling within the acceptance region $[-\lambda, \lambda]$.

If the estimate $\hat{\theta} = \bar{z}$ of θ falls outside the interval $[-\lambda, \lambda]$, then it is said to be a **significant estimate** — hypothesis H_1 is accepted because H_0 is **insignificant**.

Note that while \bar{z} is the **point estimate** of θ , the confidence region around \bar{z} can be seen as an **interval estimate** of the parameter θ .

Figure 1.5.2-2 illustrates these concepts.

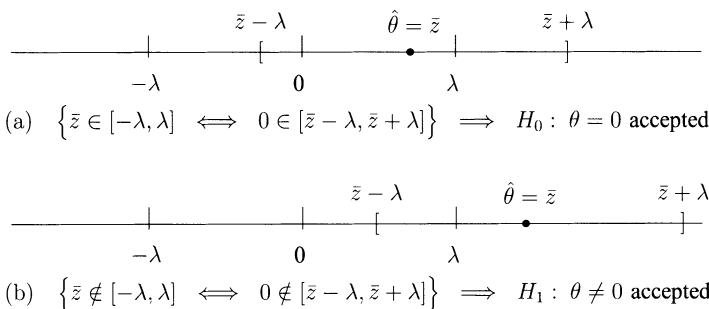


Figure 1.5.2-2: Parameter estimate. (a) $\hat{\theta}$ insignificant; (b) $\hat{\theta}$ significant.

Thus the estimate of a parameter that might be zero is accepted as **significant** (at a “level” of 5% for the null hypothesis) if

$$\frac{|\hat{\theta}|}{\sigma_{\hat{\theta}}} > 1.96 \quad (1.5.2-17)$$

where the threshold has been determined under the normal assumption and

$$\sigma_{\hat{\theta}} = \frac{\sigma}{\sqrt{n}} \quad (1.5.2-18)$$

is the standard deviation of the estimate, or its **standard error**.

Defining the **estimation error** as

$$\tilde{\theta} \triangleq \hat{\theta} - \theta \quad (1.5.2-19)$$

one can see that

$$\begin{aligned} p(\hat{\theta}|\theta) &= \mathcal{N}(\hat{\theta}; \theta, \sigma_{\hat{\theta}}^2) = \mathcal{N}(\hat{\theta} - \theta; 0, \sigma_{\hat{\theta}}^2) \\ &= \mathcal{N}(\tilde{\theta}; 0, \sigma_{\hat{\theta}}^2) \triangleq p(\tilde{\theta}) = \mathcal{N}(\tilde{\theta}; 0, \sigma_{\tilde{\theta}}^2) \end{aligned} \quad (1.5.2-20)$$

Therefore (1.5.2-18), which is the **standard deviation of the estimate** (about the true mean), $\sigma_{\hat{\theta}}$, is the same as $\sigma_{\tilde{\theta}}$, the **standard deviation of the estimation error** (which is zero mean), that is,

$$\sigma_{\hat{\theta}} = \sigma_{\tilde{\theta}} \quad (1.5.2-21)$$

The (statistical) significance of parameter estimates is used in choosing the order of models to describe system equations (e.g., for target motion). (See also problem 1-2.)

Example

A parameter θ is estimated based on n measurements, each with an additive error with standard deviation $\sigma = 10$. Assume the estimate obtained from $n = 10$ measurements is $\hat{\theta} = 2$. The significance statistic (at a level of 5%) is in this case

$$\frac{|\hat{\theta}|}{\sigma_{\hat{\theta}}} = \frac{|\hat{\theta}|}{\sigma/\sqrt{n}} = \frac{2}{10/\sqrt{10}} = 0.63 \ll 1.96 \quad (1.5.2-22)$$

i.e., based on these $n = 10$ measurements, this estimate is clearly not statistically significant — the true value of the parameter can be assumed as zero. The threshold above corresponds to 5% level of significance (5% probability that the estimate is “due to noise only,” or 95% confidence that it is nonzero).

Assume next that the observations are repeated until one has $n = 100$ measurements and that the estimate is still $\hat{\theta} = 2$. The significance statistic is now

$$\frac{|\hat{\theta}|}{\sigma/\sqrt{n}} = \frac{2}{10/\sqrt{100}} = 2 > 1.96 \quad (1.5.2-23)$$

i.e., based on these $n = 100$ measurements, this estimate is statistically significant, even though only marginally.

Note that it is somewhat arbitrary to use a 5% level of significance level for the threshold for acceptance/rejection of the null hypothesis. A marginally significant parameter can become not significant if the above level is lowered. However, while there is some subjectivity in the selection of the level for the test, this only affects marginal cases. For example, if $\hat{\theta} = 12$, then this value would be accepted even from $n = 10$ measurements since the test statistic is

$$\frac{|\hat{\theta}|}{\sigma/\sqrt{n}} = \frac{12}{10/\sqrt{10}} = 3.8 \quad (1.5.2-24)$$

i.e., large enough to exceed even the 1% level threshold, which is 3.38.

1.5.3 Monte Carlo Runs and Comparison of Algorithms

The performance of an estimation (or control) algorithm is usually quantified by the expected value of a “cost” function

$$J = E[C] \quad (1.5.3-1)$$

If C is, for instance, the squared error in the estimation of a certain variable, then J is the corresponding mean square error.

In many situations, the performance of an algorithm of interest cannot be evaluated analytically.¹³ In such a case, **Monte Carlo simulations (runs)** are made to obtain an estimate of J from a sample average of *independent* realizations C_i , $i = 1, \dots, N$, of the cost C . The larger the number of such runs, the smaller is the variability (error) of the resulting estimate. Also, a larger number of runs increases the power of the hypothesis testing used in comparing different algorithms.

The estimate of the performance from N independent runs is the **sample average** (or **sample mean**) of the N realizations of the cost

$$\bar{C} = \frac{1}{N} \sum_{i=1}^N C_i \quad (1.5.3-2)$$

with the associated standard deviation — the **standard error**

$$\sigma_{\bar{C}} = \sqrt{\frac{1}{N^2} \sum_{i=1}^N (C_i - \bar{C})^2} \quad (1.5.3-3)$$

The above follows from the fact that the variance of the sample mean \bar{C} is the variance of C divided by N ; since this variance is not known, it is replaced by the sample variance, which has another N in the denominator. This is the reason for having N^2 in the denominator of (1.5.3-3).

¹³Those who can, do. Those who cannot, simulate. Those who cannot simulate, supervise.

These can be used to obtain a confidence region for the performance assuming that its distribution is, in view of the central limit theorem, approximately normal. Since the pdf of (1.5.3-2) is in general skewed, confidence regions based on the normal assumption might not be accurate unless N is very large — possibly of the order of thousands.

Algorithm Comparison as a Hypothesis Testing Problem

When two algorithms are compared based on simulations, this can be formulated as a hypothesis testing problem as follows.

Assume one has the sample mean of the performance (e.g., MSE, or some other cost to be minimized) for algorithm j from N independent runs

$$\bar{C}^{(j)} = \frac{1}{N} \sum_{i=1}^N C_i^{(j)} \quad j = 1, 2 \quad (1.5.3-4)$$

where $C_i^{(j)}$ is the performance of algorithm j in run i and $C_i^{(j)}$ is independent of $C_k^{(j)}$, $\forall k \neq i$. Each pair of runs i uses the *same random variables for the two algorithms* $j = 1, 2$. As shown in problem 1-3, it is *beneficial* to use the same random variables for both algorithms in the same run.

Some Simple Comparison Techniques

Assume that from the simulations one has

$$\bar{C}^{(1)} < \bar{C}^{(2)} \quad (1.5.3-5)$$

This does not necessarily imply that algorithm 1 is better than algorithm 2. Any such statement must be qualified by a probability α of error of type I.

Comparison technique #0, based on (1.5.3-5) without any statistical analysis of the individual run outcomes, is clearly *naive*.

Comparison technique #1 — the simplest (but still simplistic) statistical approach — consists of the following:

1. Calculate confidence regions around the two sample means (based on (1.5.3-2) and (1.5.3-3)).
2. If they do not overlap, declare that the algorithm with the smaller sample mean is the superior one.

The last step above implicitly assumes the two sample means to be uncorrelated. Note that this is *incorrect*, since each pair of runs uses the same random numbers.

The Optimal Comparison Technique

Since $C_i^{(1)}$ is correlated with $C_i^{(2)}$ (because they use the same random variables), the optimal test, which will be based on *independent samples*, is as follows.

The hypothesis testing problem is

$$H_0 : \Delta = J^{(2)} - J^{(1)} \leq 0 \quad (\text{algorithm 1 not better than 2}) \quad (1.5.3-6)$$

versus

$$H_1 : \Delta = J^{(2)} - J^{(1)} > 0 \quad (\text{algorithm 1 better than 2}) \quad (1.5.3-7)$$

subject to

$$P\{\text{accept } H_1 | H_0 \text{ true}\} = \alpha \quad (\text{level of significance (of hypothesis } H_0)) \quad (1.5.3-8)$$

where

$$J^{(j)} = E[C^{(j)}] \quad j = 1, 2 \quad (1.5.3-9)$$

are the *true expected values of the cost functions (true performance)*.

The decision whether H_1 (“1” better than “2”) should be accepted is made based upon the *sample performance differences*

$$\boxed{\Delta_i = C_i^{(2)} - C_i^{(1)}} \quad (1.5.3-10)$$

Note that Δ_i is independent of Δ_k , $\forall k \neq i$.

Then, H_1 is accepted if

$$\boxed{\mu \triangleq \frac{\bar{\Delta}}{\sigma_{\bar{\Delta}}} > \mu_0} \quad (1.5.3-11)$$

where

$$\boxed{\bar{\Delta} = \frac{1}{N} \sum_{i=1}^N \Delta_i} \quad (1.5.3-12)$$

and

$$\boxed{\sigma_{\bar{\Delta}} = \sqrt{\frac{1}{N^2} \sum_{i=1}^N (\Delta_i - \bar{\Delta})^2}} \quad (1.5.3-13)$$

are the *sample mean* of the differences (1.5.3-10) and the *standard error* of this sample mean, respectively.

Assuming the error in $\bar{\Delta}$ to be normal, the threshold μ_0 is based on the *upper tail* of the normal density: $\mu_0 = 1.64$ for $\alpha = 5\%$, $\mu_0 = 2.33$ for $\alpha = 1\%$, and so on (see Subsection 1.5.4). This follows from the fact that the test is for *positive mean (H_1) versus zero or negative mean* and H_1 is accepted if and only if the sample mean (1.5.3-12) is positive and statistically significant (1.5.3-11); α is the probability of accepting H_1 (positive mean) when the true mean is zero.

This procedure is summarized below:

$C_1^{(1)}$	$C_1^{(2)}$	Δ_1
\vdots	\vdots	\vdots
$C_N^{(1)}$	$C_N^{(2)}$	Δ_N
$\bar{C}^{(1)}$	$\bar{C}^{(2)}$	$\bar{\Delta}, \sigma_{\bar{\Delta}}$

Remarks

The optimal algorithm comparison is based on the significance test of the *sample mean of the differences* (the mean of the terms in the last column above) instead of the naive comparison of the performance estimates from the first two columns (technique #0) or using two confidence regions around these estimates (technique #1).

The applicability of the CLT is much more realistic on the average difference (1.5.3-12) than on the average outcomes (1.5.3-4) since the differences are much less skewed (and independent from run to run).

The fact that the same random numbers are used for the two algorithms leads to a *positive correlation* between the two sample means, and the optimal test presented takes advantage of this. (See also problem 1-3.)

Section 11.5 presents a major example of the application of this optimal comparison technique.

1.5.4 Tables of the Chi-Square and Gaussian Distributions

Table 1.5.4-1 presents the points x on the chi-square distribution for a given *upper tail probability*

$$Q = P\{y > x\} \quad (1.5.4-1)$$

where

$$y \sim \chi_n^2 \quad (1.5.4-2)$$

and n is the number of degrees of freedom. This tabulated function is also known as the *complementary distribution*.

An alternative way of writing (1.5.4-1) is

$$x(1 - Q) \triangleq \chi_n^2(1 - Q) \quad (1.5.4-3)$$

which indicates that to the left of the point x the probability mass is $1 - Q$. This is the $100(1 - Q)$ *percentile point*.

The 95% probability region for a χ_2^2 variable can be taken as the *one-sided probability region* (cutting off the 5% *upper tail*)

$$[0, \chi_2^2(0.95)] = [0, 5.99] \quad (1.5.4-4)$$

or the **two-sided probability region** (cutting off both 2.5% tails)

$$[\chi_2^2(0.025), \chi_2^2(0.975)] = [0.05, 7.38] \quad (1.5.4-5)$$

For a χ_{100}^2 variable, the two-sided 95% probability region is

$$[\chi_{100}^2(0.025), \chi_{100}^2(0.975)] = [74, 130] \quad (1.5.4-6)$$

Note the skewedness of the chi-square distribution: The above two-sided regions are not symmetric about the corresponding means (2 and 100, respectively, in the above two equations).

For degrees of freedom above 100, the following approximation of the points on the chi-square distribution can be used:

$$\chi_n^2(1 - Q) = \frac{1}{2} [\mathcal{G}(1 - Q) + \sqrt{2n - 1}]^2 \quad (1.5.4-7)$$

where \mathcal{G} is defined in (1.5.4-11). For example, $\chi_{400}^2(0.025) = 346$, $\chi_{400}^2(0.975) = 457$.

The last line of Table 1.5.4-1 shows the points x on the **standard (zero mean and unity variance) normal (Gaussian)** distribution for the same tail probabilities. In this case

$$y \sim \mathcal{N}(0, 1) \quad (1.5.4-8)$$

and, with

$$Q = P\{y > x\} \quad (1.5.4-9)$$

these points will be denoted as

$$x(1 - Q) \triangleq \mathcal{G}(1 - Q) \quad (1.5.4-10)$$

Thus, with this notation the 95% **two-sided probability region** for an $\mathcal{N}(0, 1)$ random variable is

$$[\mathcal{G}(0.025), \mathcal{G}(0.975)] = [-1.96, 1.96] \quad (1.5.4-11)$$

In terms of the cumulative distribution function (cdf) of a standard Gaussian random variable

$$\begin{aligned} P_{\mathcal{G}}(x) &\triangleq P\{y \leq x\} \\ &= \int_{-\infty}^x \mathcal{N}(y; 0, 1) dy \\ &= \frac{1}{\sqrt{2\pi}} \int_{-\infty}^x e^{-\frac{y^2}{2}} dy \end{aligned} \quad (1.5.4-12)$$

one can write (1.5.4-11) as

$$P_{\mathcal{G}}(1.96) = 0.975 \quad (1.5.4-13)$$

$$P_{\mathcal{G}}(-1.96) = 0.025 \quad (1.5.4-14)$$

Note that $P_{\mathcal{G}}$ defined in (1.5.4-12) is the inverse function of \mathcal{G} defined in (1.5.4-10), that is, \mathcal{G} is the inverse cdf.

The 95% **one-sided probability region** for a standard Gaussian random variable is given by its 95% point:

$$(-\infty, \mathcal{G}(0.95)] = (-\infty, 1.64] \quad (1.5.4-15)$$

Table 1.5.4-1: Tail probabilities of the chi-square and normal densities.

$n \setminus Q$	0.99	0.975	0.95	0.90	0.75	0.5	0.25	0.10	0.05	0.025	0.01	5E-3	1E-3
1	2E-4	.001	.003	.016	.102	.455	1.32	2.71	3.84	5.02	6.63	7.88	10.8
2	.020	.051	.103	.211	.575	1.39	2.77	4.61	5.99	7.38	9.21	10.6	13.8
3	.115	.216	.352	.584	1.21	2.37	4.11	6.25	7.81	9.35	11.3	12.8	16.3
4	.297	.484	.711	1.06	1.92	3.36	5.39	7.78	9.49	11.1	13.3	14.9	18.5
5	.554	.831	1.15	1.61	2.67	4.35	6.63	9.24	11.1	12.8	15.1	16.7	20.5
6	.872	1.24	1.64	2.20	3.35	5.35	7.84	10.6	12.6	14.4	16.8	18.5	22.5
7	1.24	1.69	2.17	2.83	4.25	6.35	9.04	12.0	14.1	16.1	18.5	20.3	24.3
8	1.65	2.18	2.73	3.49	5.07	7.34	10.2	13.4	15.5	17.5	20.1	22.0	26.1
9	2.09	2.70	3.33	4.17	5.90	8.34	11.4	14.7	17.0	19.0	21.7	23.6	27.9
10	2.56	3.25	3.94	4.87	6.74	9.34	12.5	16.0	18.3	20.5	23.2	25.2	29.6
11	3.05	3.82	4.57	5.58	7.58	10.3	13.7	17.3	19.7	22.0	24.7	26.8	31.3
12	3.57	4.40	5.23	6.30	8.44	11.3	14.8	18.5	21.0	23.3	26.2	28.3	32.9
13	4.11	5.01	5.90	7.04	9.30	12.3	16.0	19.8	22.4	24.7	27.7	29.8	34.5
14	4.66	5.63	6.57	7.79	10.2	13.3	17.1	21.1	23.7	26.1	29.1	31.3	36.1
15	5.23	6.26	7.26	8.55	11.0	14.3	18.2	22.3	25.0	27.5	30.6	32.8	37.7
16	5.81	6.91	7.96	9.31	11.9	15.3	19.4	23.5	26.3	28.8	32.0	34.3	39.3
17	6.41	7.56	8.67	10.1	12.8	16.3	20.5	24.8	27.6	30.2	33.4	35.7	40.8
18	7.01	8.23	9.40	10.9	13.7	17.3	21.6	26.0	28.9	31.5	34.8	37.2	42.3
19	7.63	8.91	10.1	11.7	14.6	18.3	22.7	27.2	30.1	32.9	36.2	38.6	43.8
20	8.26	9.60	10.9	12.4	15.5	19.3	23.8	28.4	31.4	34.2	37.6	40.0	45.3
25	11.5	13.1	14.6	16.5	19.9	24.3	29.3	34.4	37.7	40.6	44.3	46.9	52.6
30	15.0	16.8	18.5	20.6	24.5	29.3	34.8	40.3	43.8	47.0	50.9	53.7	59.7
40	22.2	24.4	26.5	29.1	33.7	39.3	45.6	51.8	55.8	59.3	63.7	66.8	73.4
50	29.7	32.4	34.8	37.7	43.0	49.3	56.3	63.2	67.5	71.4	76.2	79.5	86.7
60	37.5	40.5	43.2	46.5	52.3	59.3	67.0	74.4	79.1	83.3	88.4	92.0	99.6
70	45.4	48.8	51.7	55.3	61.7	69.3	77.6	85.5	90.5	95.0	100	104	112
80	53.5	57.2	60.4	64.2	71.1	79.3	88.1	96.6	102	107	112	116	125
90	61.8	65.6	69.1	73.3	80.6	89.3	98.6	108	113	118	124	128	137
100	70.1	74.2	77.9	82.4	90.1	99.3	109	118	124	130	136	140	149
150	112	118	123	128	138	149	161	173	180	186	193	198	209
200	157	163	168	175	186	199	213	226	234	241	249	255	268
300	246	254	261	269	283	299	316	332	341	350	360	367	381
400	337	346	355	364	381	399	419	437	448	457	469	478	495
500	429	440	449	460	478	499	521	541	553	563	576	585	603
800	710	724	735	749	773	799	827	852	867	880	896	907	929
$\bar{G}(0)$	-2.33	-1.96	-1.64	-1.28	-0.675	0	.675	1.28	1.64	1.96	2.33	2.58	3.09

Some Distribution-free Probability Regions

The *Chebyshev inequality* provides a *distribution-free probability region*. For a random variable x with mean μ and variance σ^2 , one has, for any $\epsilon > 0$,

$$P\{|x - \mu| \geq \epsilon\} \leq \frac{\sigma^2}{\epsilon^2} \quad (1.5.4-16)$$

regardless of the distribution of x .

For a non-negative random variable x (i.e., $p(x) = 0 \ \forall x < 0$) with mean μ , the *Markov inequality* states that, for any $\alpha > 0$,

$$P\{x \geq \alpha\} \leq \frac{\mu}{\alpha} \quad (1.5.4-17)$$

1.6 NOTES AND PROBLEMS

1.6.1 Bibliographical Notes

An excellent historical survey of the developments of estimation theory, starting from the least squares technique of Gauss and Legendre, the maximum likelihood technique of R. A. Fisher, the Wiener filtering for signals in noise, and the emergence of digital computer oriented recursive algorithms developed by Swerling, Kalman, Bucy, and others, can be found in [Sorenson85].

Other texts on the same topic as this one are [Meditch69, Sage71, Gelb74, Anderson79, Lewis86, Maybeck79, Maybeck82, Grewal93, Brown97].

For background material, further references on linear algebra as well as linear systems are [Zadeh63, Fortmann77, Chen84]. The existing literature on this topic consists of tens of books, and most of them are equally good. A succinct summary of this material can be found in [Gelb74]. The probabilistic tools are covered in [Papoulis84] and succinctly in [Gelb74]. Spectral factorization is discussed in [Poor88]. The statistical tools (hypothesis testing) can be found, for example, in [Melsa78, Johnston72], among others.

For the more specialized target tracking topics — in particular, tracking with uncertain origin observations and multitarget tracking — the reader is referred to [Blackman86, Bar-Shalom88]. The texts [Bar-Shalom95] and [Blackman99] are the most recent compilations of algorithms in this area, with suitable explanations and design guidelines.

1.6.2 Problems

1-1 Independence versus conditional independence. Consider x and y independent Gaussian random variables with zero mean and unity variance.

1. For an arbitrary new random variable does the following hold

$$p(x, y|z) = p(x|z)p(y|z) \quad ?$$

2. Let $z = x + y$. Write the explicit expression of the pdf $p(x, y|z)$. Use the Dirac delta function if necessary.

- 1-2 Monte Carlo runs for low probability events.** Given an experiment with binary outcome x , assume your colleague did a theoretical calculation that indicates that

$$P\{x = 1\} = \hat{p} = 10^{-4}$$

However, you do not believe this colleague and plan to conduct N Monte-Carlo (i.i.d.) repetitions of the experiment. The outcomes are x_i , $i = 1, \dots, N$. Estimate, using a (say, 95%) confidence region, the number N of experiments you have to carry out to confirm or refute your colleague's theoretical calculation with, e.g., 10% tolerance.

- 1-3 Monte Carlo runs for algorithm comparison.** Two estimation algorithms are to be compared based on Monte Carlo runs. Assume their average performances (MSE) were obtained as $\bar{C}^{(1)} = 5$ and $\bar{C}^{(2)} = 8$ from $N = 100$ runs.

1. Can one say that algorithm 1 is superior to algorithm 2?
2. Assume the standard deviations associated with these sample means are

$$\sigma_{\bar{C}^{(j)}} = \left[\frac{1}{N^2} \sum_{i=1}^N \left[C_i^{(j)} - \bar{C}^{(j)} \right]^2 \right]^{1/2} = 1 \quad j = 1, 2$$

Using the above information answer question 1.

3. Denote the correlation coefficient between the outcomes of the runs with the two algorithms as ρ , that is,

$$E \left[(C_i^{(1)} - J^{(1)}) (C_i^{(2)} - J^{(2)}) \right] = \rho \sigma_{C^{(1)}} \sigma_{C^{(2)}}$$

Show that if $\rho = 0.5$ and the variances are as in item 2, the optimal statistical test will yield a conclusive answer.

- 1-4 Covariance of a mixture.** Prove that an equivalent expression of the "spread of the means" term (1.4.16-9) in the covariance of a mixture is

$$\tilde{P} = \sum_i \sum_{j < i} p_i p_j (\bar{x}_i - \bar{x}_j)(\bar{x}_i - \bar{x}_j)'$$

where the summation is over all the pairs of indices without repetition.

1-5 Dimension of a pdf.

1. Find the physical dimension of the pdf of the random vector x when x is
 - a. The position of a point in an n -dimensional Euclidean space.
 - b. The state of a constant (unknown) acceleration point moving along, say, the ξ -axis.
2. Find the units of the pdf of a velocity measured in furlongs/fortnight.
3. Can one compare the pdf of a random variable whose dimension is length to the joint pdf of two such random variables?

- 1-6 Log-normal random variable.** Given the random variable $y \sim \mathcal{N}(\mu, \sigma^2)$, then $x \triangleq e^y$ is called a log-normal random variable. Find its mean and variance.

- 1-7 Moments of linear transformation of random variables.** Given the random variables x and y of dimensions n_x and n_y , with means \bar{x} and \bar{y} , respectively, and with covariances P_{xx}, P_{yy}, P_{xy} :

- Find the mean and covariance of the n_z -dimensional vector

$$z = Ax + By + c$$

where A , and B are matrices of appropriate dimensions.

- Indicate the dimensions of A , B , and c .

- 1-8 Chapman-Kolmogorov equation.** Prove that the following equation holds for a (discrete time) Markov process

$$\int p[x(k)|x(k-1)]p[x(k-1)|x(k-2)] dx(k-1) = p[x(k)|x(k-2)]$$

- 1-9 Hypothesis testing with correlated noise.** Consider the hypotheses

$$H_0 : \theta = 0 \quad H_1 : \theta \neq 0$$

and the observations

$$z_i = \theta + w_i \quad i = 1, \dots, n$$

with w_i zero-mean jointly Gaussian but not independent. Denoting

$$w = [w_1 \cdots w_n]'$$

one has the covariance matrix (assumed given)

$$E[ww'] = P$$

For the above

- Specify the optimal hypothesis test for false alarm probability α .
- Solve explicitly for $n = 2$, $P = \begin{bmatrix} 1 & 0.5 \\ 0.5 & 1 \end{bmatrix}$ and $\alpha = 1\%$.

- 1-10 Partial derivative with respect to a matrix.** The partial derivative of a scalar q with respect to the matrix $A = [a_{ij}]$ is defined as

$$\frac{\partial q}{\partial A} \triangleq \left[\frac{\partial q}{\partial a_{ij}} \right]$$

Prove that

- For B symmetric, $\frac{\partial}{\partial A} \text{tr}[ABA'] = 2AB$.
- For B not symmetric, $\frac{\partial}{\partial A} \text{tr}[AB] = B'$.

- 1-11 Fourth joint moment of Gaussian random variables.** Show that, if the zero-mean scalar random variables x_i , $i = 1, \dots, 4$, are jointly Gaussian, then

$$E[x_1x_2x_3x_4] = E[x_1x_2]E[x_3x_4] + E[x_1x_3]E[x_2x_4] + E[x_1x_4]E[x_2x_3]$$

Hint: Use the characteristic function (1.4.15-3) and take a suitable fourth derivative of it.

1-12 Wiener process increments.

- Using (1.4.20-3) show that

$$E[(dw(t))^2] = \alpha dt \quad \text{and} \quad dw(t) \sim \mathcal{N}(0, \alpha dt)$$

- Using the above, show that the limit of $\frac{dw(t)}{dt}$ tends to ∞ . Hint: Find the “order of magnitude” of $dw(t)$ in terms of dt (a zero-mean random variable is of the order of its standard deviation).

1-13 Conditional pdf of the sum of two Gaussian random variables. Find the pdf (1.4.13-12) in the general case (1.4.13-6).**1-14 Probability matrix.** Find $\sum_{j=1}^n \pi_{ij}$ for (1.4.22-9).**1-15 Laplacian pdf.** Given the random variable x with Laplacian (two-sided exponential) pdf

$$p(x) = c e^{-a|x|} \quad a > 0 \quad -\infty < x < \infty$$

- Find the normalization constant c , the mean and the variance of x .
- Repeat the above if it is known that $|x| \leq b$.

1-16 Moments of a quadratic form with non-zero-mean random variables. Consider the random variables x and y with means \bar{x} and \bar{y} , respectively, and with covariances P_{xx}, P_{yy}, P_{xy} . Evaluate $E[x'Ay]$.**1-17 The law of perversity of inanimate objects.** A buttered toast falls off a table. According to the “law of perversity of inanimate objects,” the probability of it falling on the floor with the buttered side down is greater than that with the buttered side up.

The motion (rotation and fall) of the toast can be modeled as follows.

Starting with the buttered side facing up, it acquires a rotational motion with a random angular rate $\dot{\theta} \sim U[\dot{\theta}_1, \dot{\theta}_2]$, assumed to stay constant until it hits the floor. The side which will hit the floor depends on the angle it rotated.

- Indicate the interval(s) of the rotation angle θ for which the buttered side will hit the floor.
- Assuming a constant acceleration fall from the typical height of a table, h , with acceleration $g = 10 \text{ m/s}^2$, find the time the toast has to rotate.
- Find the pdf of the angle θ it rotated until it hit the floor.
- Give the expression of the probability that the buttered side will be facing down when hitting the floor.
- Evaluate the above for $h = 0.8 \text{ m}$, $\dot{\theta}_1 = 200^\circ/\text{s}$, $\dot{\theta}_2 = 300^\circ/\text{s}$.
- Reevaluate the above for the land of the Brobdingnagians, in which the size of people (and their furniture) is scaled up by a factor of 9 compared to the above.